

# Multicore Architecture Programming – Introduction

Olivier Aumage

Inria & LaBRI lab.

`olivier.aumage @ inria.fr`

2024 – 2025



# Introduction

# Context

## Rapidly evolving processor technologies

- **Density**
  - 2024
    - Intel Emerald Rapids architecture: up to 64 cores [\[link\]](#) (Dec. 2023)
      - 2024: Granite Rapids (up to 128 P-cores: HPC) / Sierra Forrest (up to 288 E-cores: Cloud Computing)
    - AMD EPYC Turin / Zen 5 architecture: up to 192 cores [\[link\]](#)
- **Complexity & Optimization trade-off**
  - Bandwidth
  - Latency
  - Energy consumption & power envelope
  - Cost
- **Heterogeneity**
  - Specialized instruction sets
  - Accelerators & specialized processors

# Objectives

## Purpose of the course

- Understand key design principles of modern processors
- Learn and combine techniques to write efficient programs for such processors
- Get ready to adapt in a rapidly evolving landscape

# Organisation

## Speakers

- **Olivier Aumage**
  - General purpose multicore programming
  - Single Instruction Multiple Data (SIMD) programming
  - Task parallelism
- **Amina Guermouche**
  - Energy & power consumption in multicore processors
  - Accelerator architectures programming

# Organisation

## Evaluation

- **Assignment + Oral defense**
  - Apply optimization techniques on sample programs

# Basic Architectural Notions


Refresh

# Common Architectural Principles of Microprocessors

## Schematic Design

- **Die areas**

- Control unit
  - Register file
- Arithmetic & logic unit (ALU)
  - Integer operations
- Floating point unit (FPU)
- Cache (s)
  - Purposes
  - Levels






# Common Architectural Principles of Microprocessors

## Schematic Design

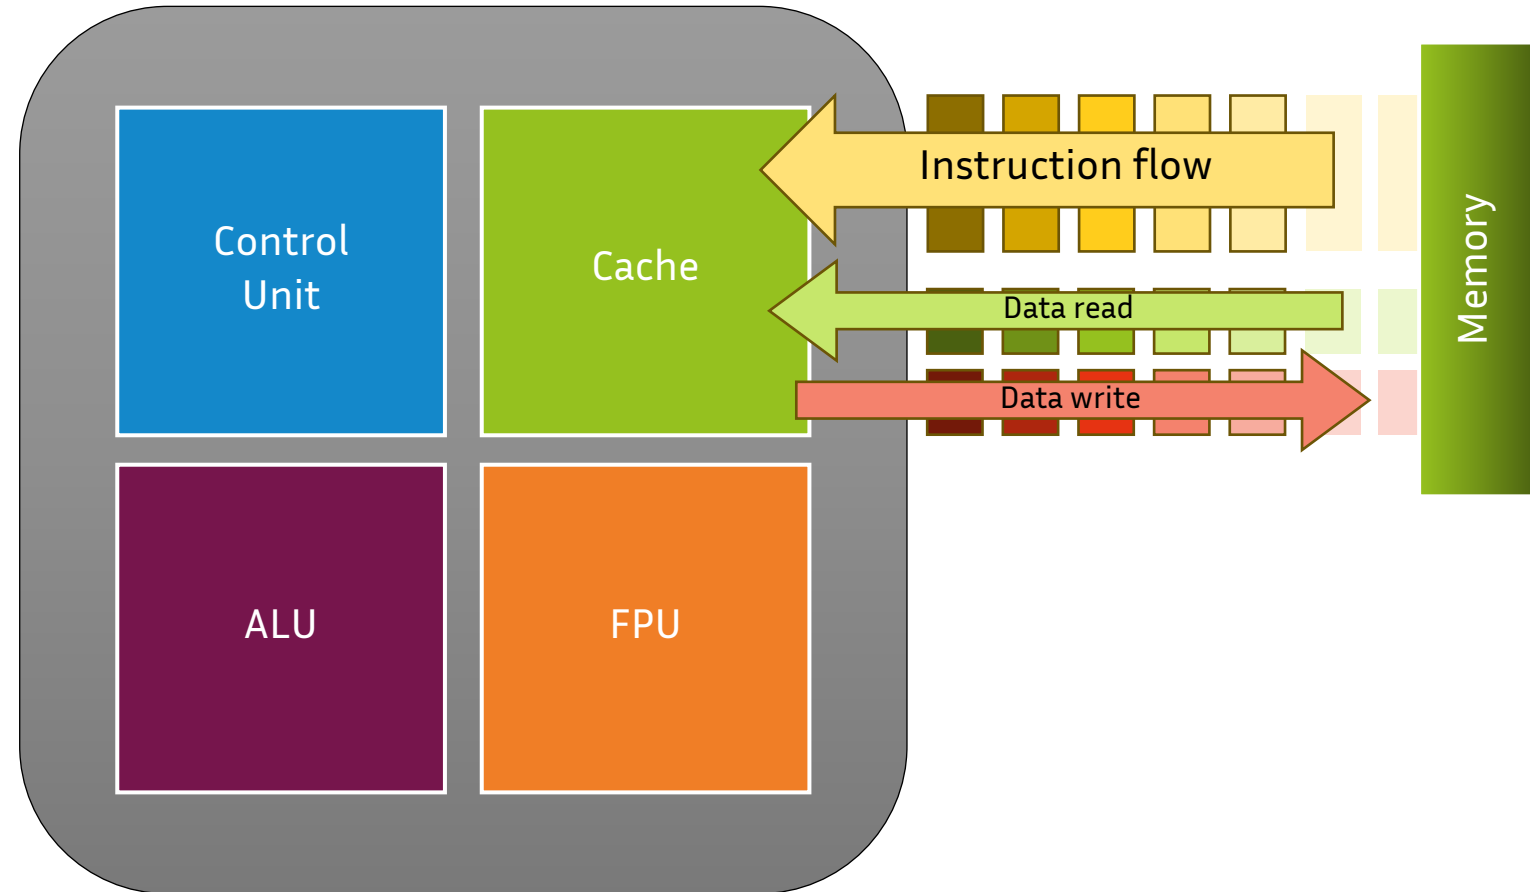
- **Die areas**
  - Control unit
    - Register file
  - Execution units
    - Arithmetic & logic unit (ALU)
    - Floating point unit (FPU)
  - Cache (s)
    - Purposes
    - Levels
- **Instruction flow**
  - Thread



# Common Architectural Principles of Microprocessors

## Schematic Design

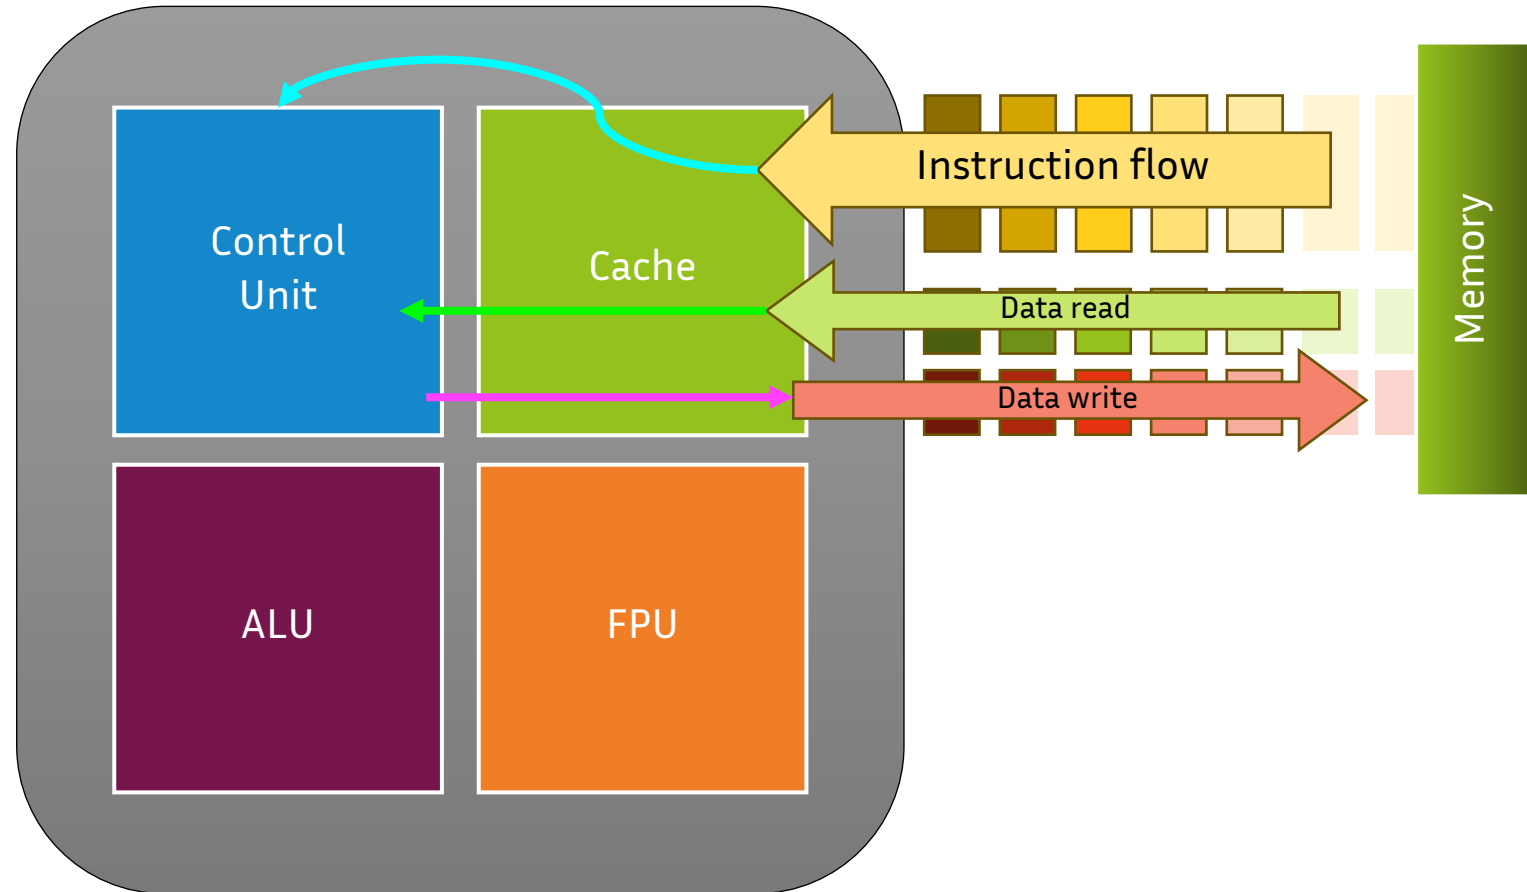
- **Die areas**
  - Control unit
    - Register file
  - Execution units
    - Arithmetic & logic unit (ALU)
    - Floating point unit (FPU)
  - Cache (s)
    - Purposes
    - Levels
- **Instruction flow**
  - Thread



# Common Architectural Principles of Microprocessors

## Schematic Design


- **Die areas**
  - Control unit
    - Register file
  - Execution units
    - Arithmetic & logic unit (ALU)
    - Floating point unit (FPU)
  - Cache (s)
    - Purposes
    - Levels
- **Instruction flow**
  - Thread



# Improving Performances

## Computing faster?


- Increase frequency?



# Frequency Wall

## Physical limits


- Thermal density
- Processor cooling



# Heat Dissipation Barrier

## Causes & Consequences

- Processor frequency evolution



Source: K. Rupp [\[link\]](#)

# Moore's Law

Gordon Moore, Intel


The number of transistors in an IC doubles every two years

## Moore's Law: The number of transistors on microchips doubles every two years

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important for other aspects of technological progress in computing – such as processing speed or the price of computers.

Our World  
in Data

### Transistor count



Data source: Wikipedia ([wikipedia.org/wiki/Transistor\\_count](https://wikipedia.org/wiki/Transistor_count))


OurWorldinData.org – Research and data to make progress against the world's largest problems.

Licensed under CC-BY by the authors Hannah Ritchie and Max Roser.

Source: Max Roser, Hannah Ritchie, 2020

Our World  
in Data

The number of people in an 10 years



Licensed under [CC-BY](#) by the authors Hannah Ritchie and Max Roser.


Team STORM – Inria &amp; LaBRI



# Dennard Scaling


Robert H. Dennard, IBM

- As transistors gets smaller, their power density stays constant
  - Power use in proportion with area
- Transition around 2006 – 2008



Source: M. Feldman, The Next Platform [\[link\]](#)

# PERFORMANCE DEVELOPMENT



Source: M. Feldman, The Next Platform [\[link\]](#)


Team STORM — Inria & LaBRI

# Heat Dissipation Barrier

## Causes & Consequences

- Processor frequency evolution
- Moore's Law
- Dennard Scaling

Improve performances by other means



Source: K. Rupp [\[link\]](#)


# Heat Dissipation Barrier

## Causes & Consequences

- Processor frequency evolution
- Moore's Law
- Dennard Scaling

## Improve performances by other means

- Doing more by unit of time
- **Parallelism**




Source: K. Rupp [\[link\]](#)

# Doing More by Unit of Time

## Hardware parallelism


- **Instruction-level parallelism (ILP)**
  - Parallelism within a single instruction flow
- **Thread-level parallelism (TLP)**
  - Parallelism across multiple instruction flows



# Doing More by Unit of Time

## Hardware parallelism

- **Instruction-level parallelism (ILP)**
  - Parallelism within a single instruction flow
- **Thread-level parallelism (TLP)**
  - Parallelism across multiple instruction flows



# Instruction-Level Parallelism

**Extract parallelism within a single flow of instructions**

- **Overlap the processing of successive instructions**
  - Pipelining
- **Process several instructions at a given time**
  - Overlapping / offload
  - Superscalar processing
- **Apply an instruction on multiple data**
  - SIMD instruction sets

# Instruction-Level Parallelism

Extract parallelism within a single flow of instructions


- **Overlap the processing of successive instructions**
  - **Pipelining**
- **Process several instructions at a given time**
  - Overlapping / offload
  - Superscalar processing
- **Apply an instruction on multiple data**
  - SIMD instruction sets



# Pipelining

## Overlap Instruction Processing Steps

- Multiple instruction execution steps
- 1 step  $\geq$  1 clock cycle

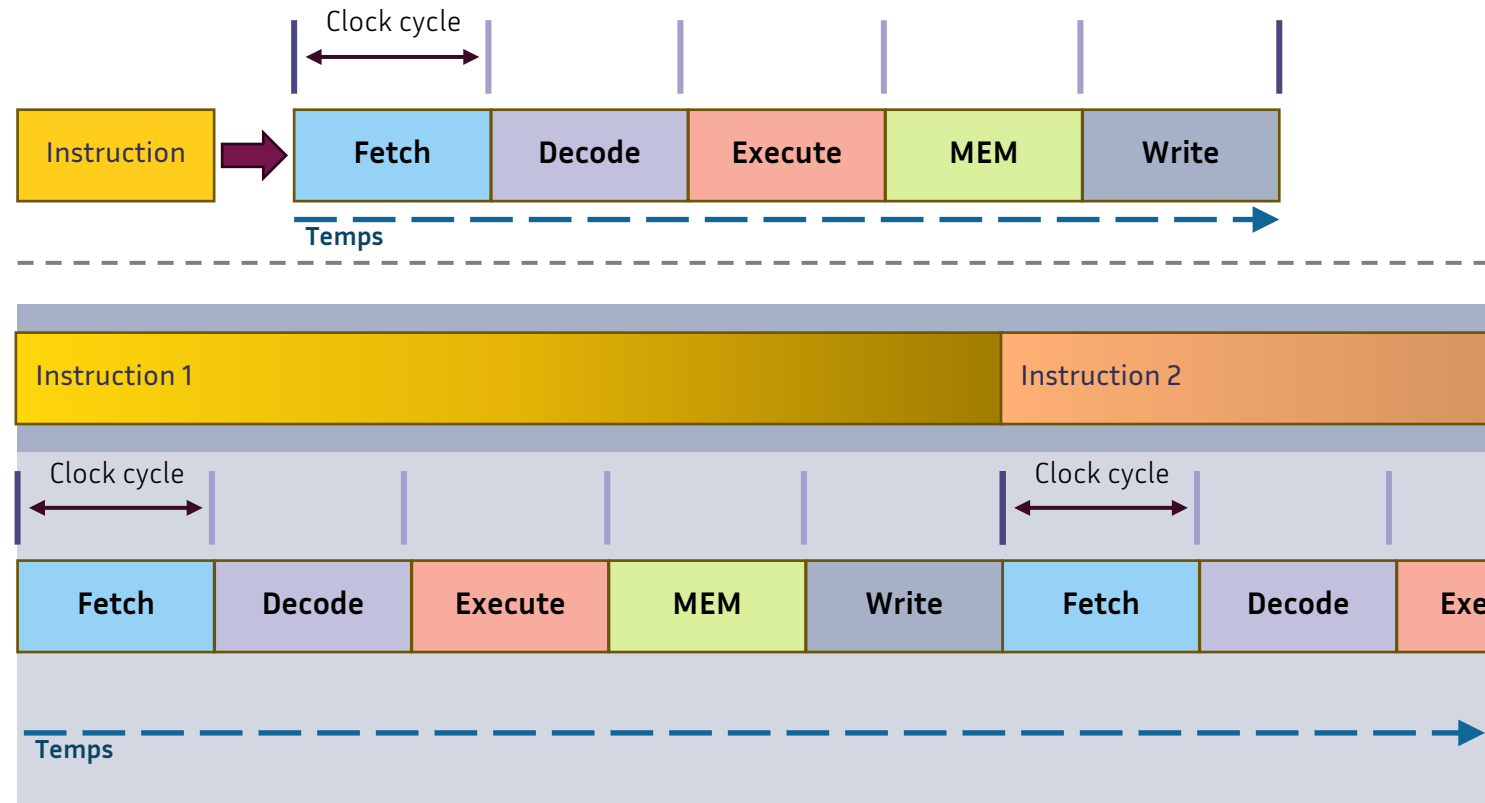


Typical RISC instruction processing sequence

# Pipelining

## Overlap Instruction Processing Steps

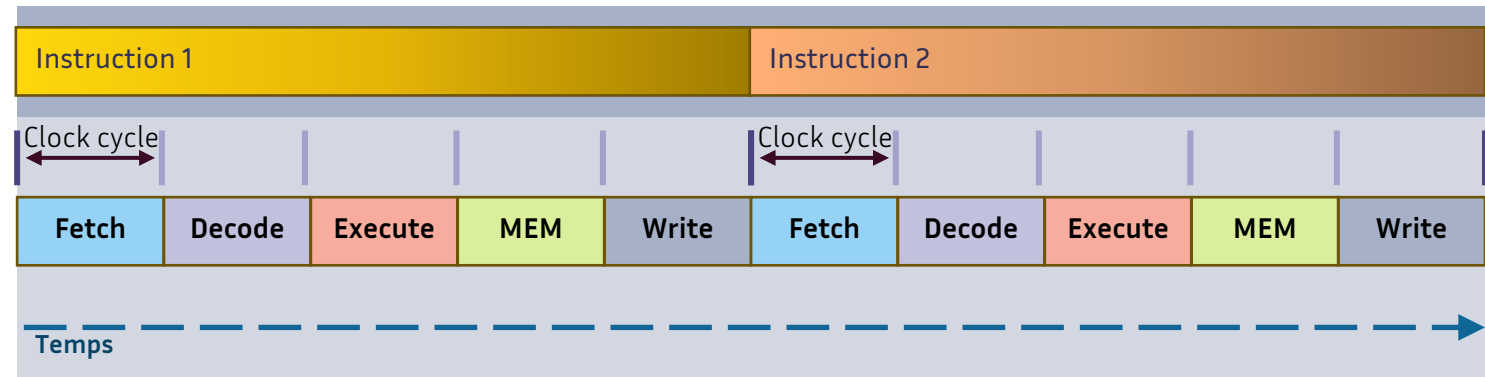
- Multiple instruction execution steps
- 1 step  $\geq$  1 clock cycle



# Pipelining

## Overlap Instruction Processing Steps

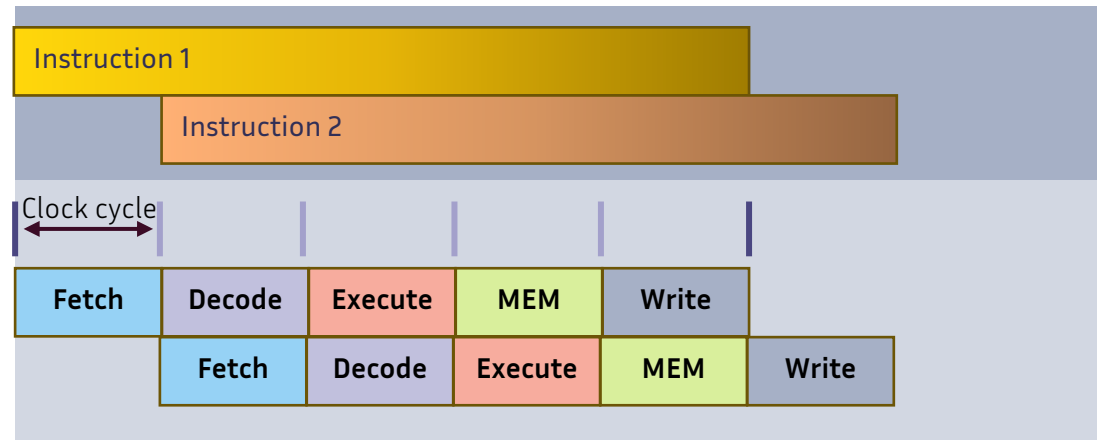
- Multiple instruction execution steps
- 1 step  $\geq$  1 clock cycle



# Pipelining

## Overlap Instruction Processing Steps

- Multiple instruction execution steps
- 1 step  $\geq$  1 clock cycle
- **Pipelining**
  - Process distinct steps of successive instructions in parallel
- **Improve Instruction per Cycle rate (IPC)**
  - Ideally, 1 instruction per cycle
- **Cost**
  - Additional circuitry to manage multiple steps in parallel



# Pipelining

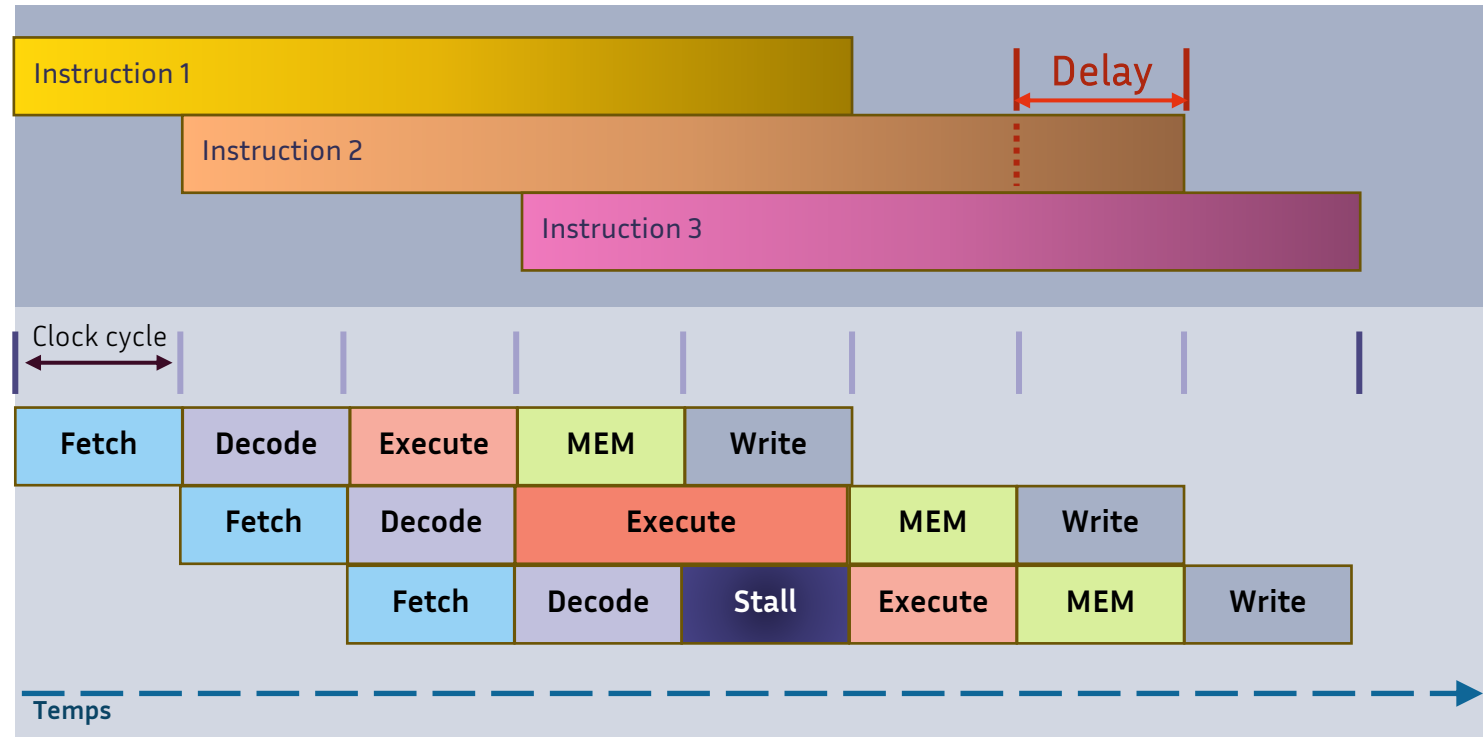
## Overlap Instruction Processing Steps

- **Limitations**

- Pipeline stalls (bubbles)

- **Causes**

- Long instructions
- Memory accesses
- Data dependencies



# Instruction-Level Parallelism

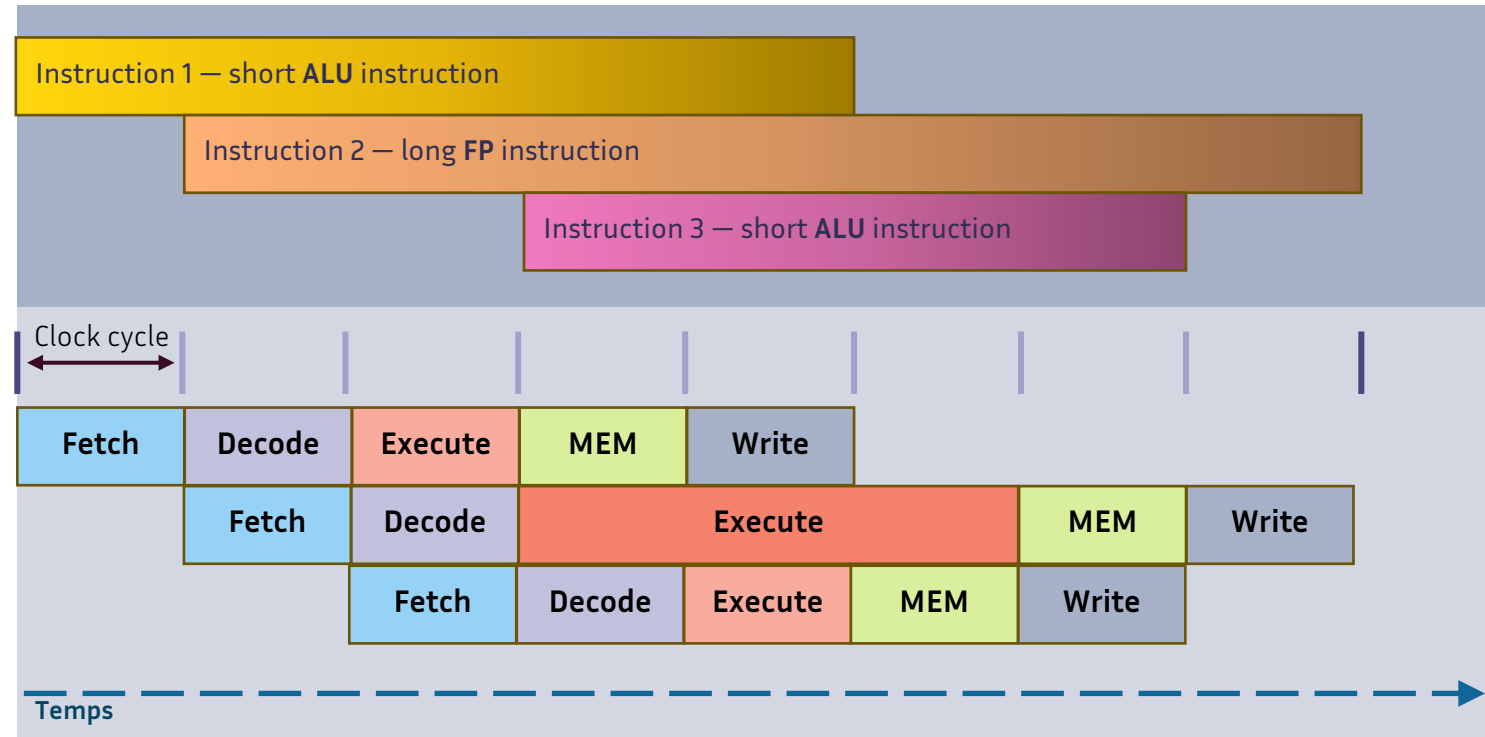
Extract parallelism within a single flow of instructions

- **Overlap the processing of successive instructions**
  - Pipelining
- **Process several instructions at a given time**
  - Overlapping / offload
  - Superscalar processing
- **Apply an instruction on multiple data**
  - SIMD instruction sets

# Overlapping / offload

## Overlap instructions processing

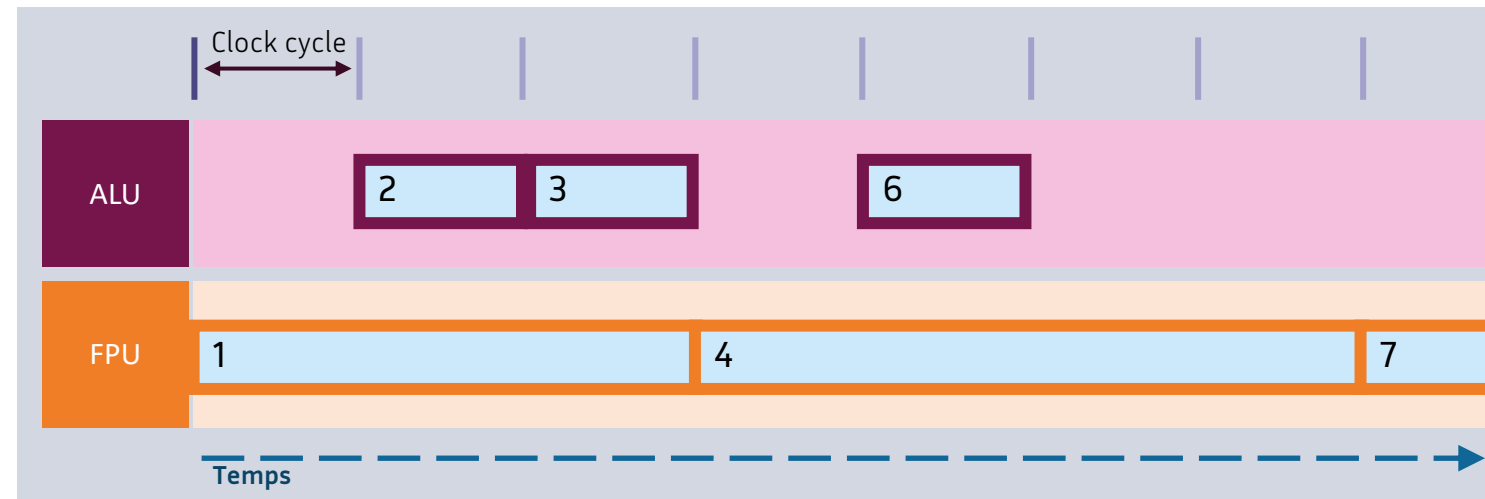
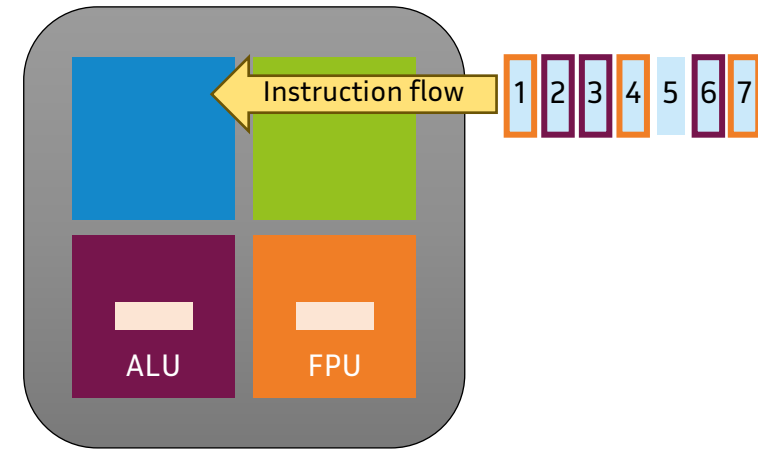
- **Heterogeneous instructions**
  - Using different arithmetic units



# Overlapping / offload

## Overlap instructions processing

- **Heterogeneous instructions**
  - Using different arithmetic units






# Superscalar Processors

## Overlap instructions processing

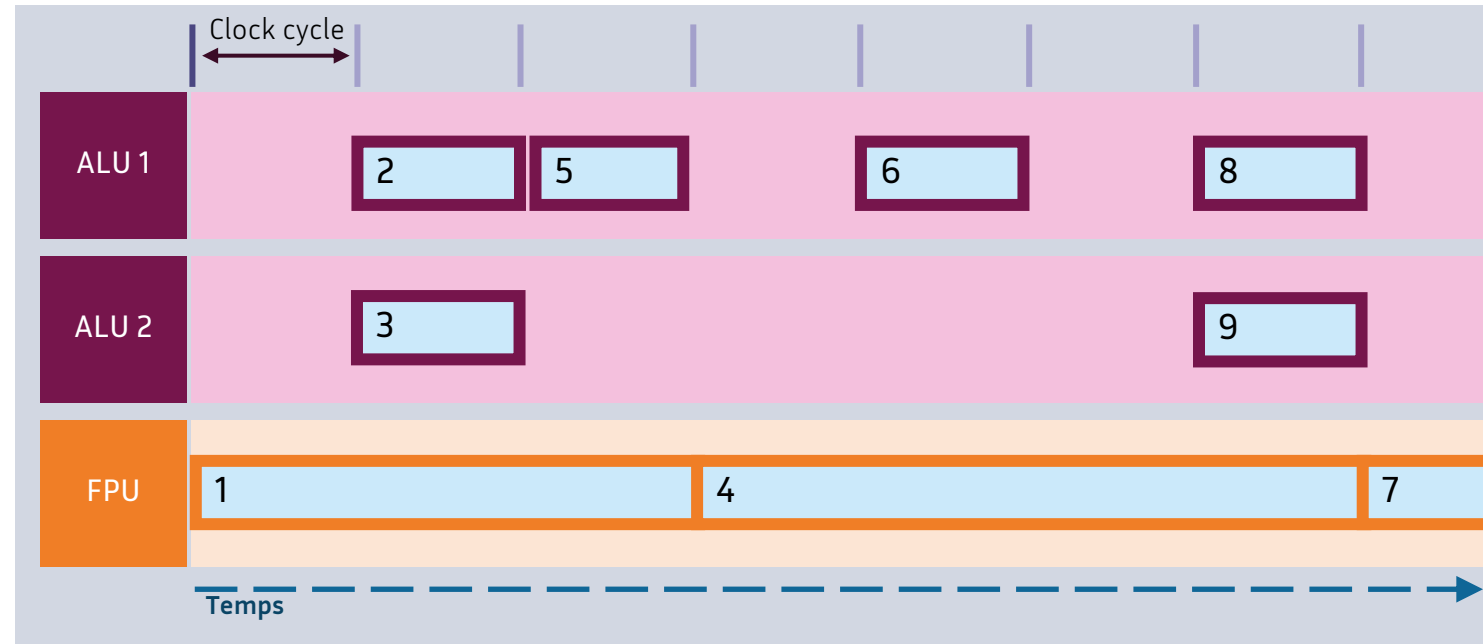
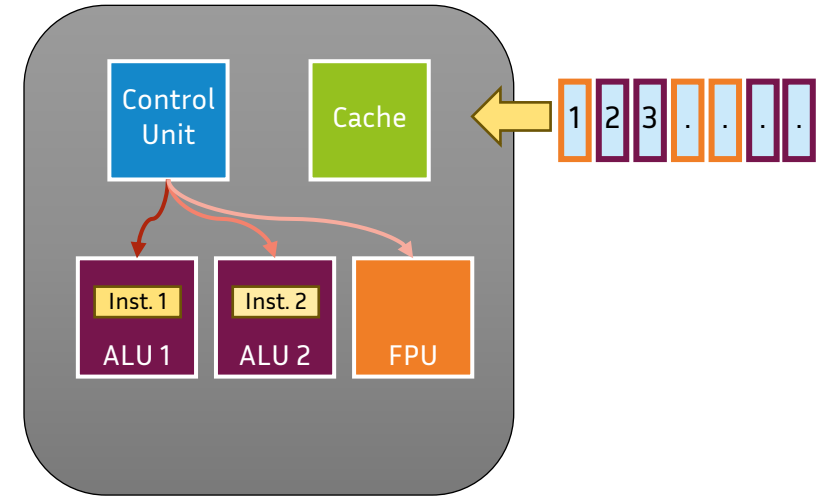
- **Homogeneous instructions**
  - Multiple identical arithmetic units



# Superscalar Processors

## Overlap instructions processing

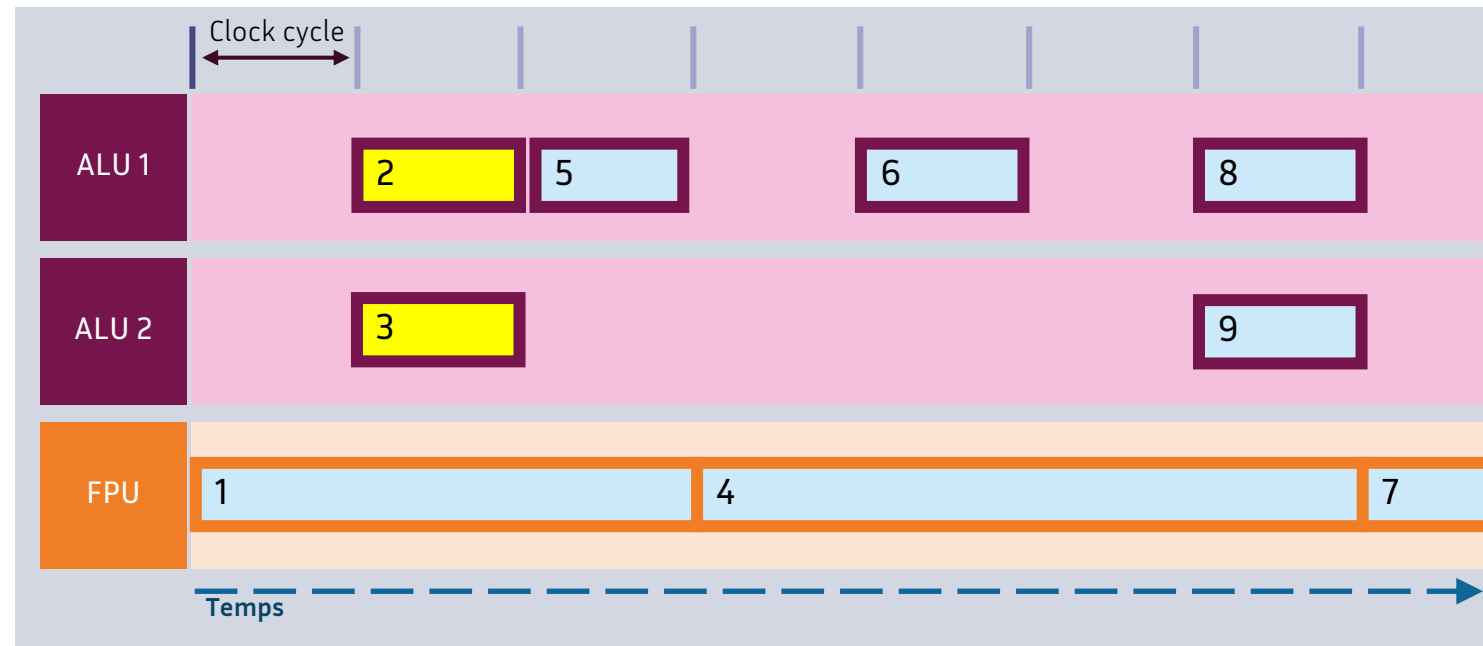
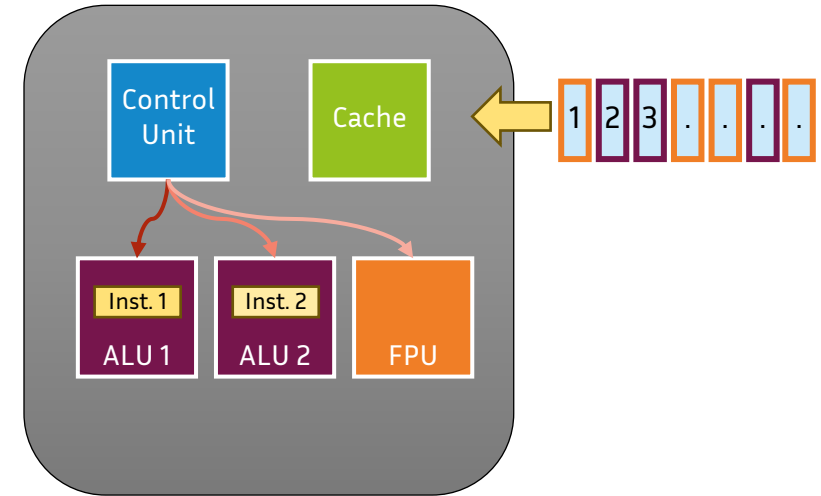
- **Homogeneous instructions**
  - Multiple identical arithmetic units
  - Instruction re-ordering



# Superscalar Processors

## Overlap instructions processing



- **Homogeneous instructions**
  - Multiple identical arithmetic units
  - Instruction re-ordering



# Superscalar Processors

## Overlap instructions processing

- **Homogeneous instructions**
  - Multiple identical arithmetic units
  - Instruction re-ordering
- **Consistency**
  - Data dependencies



# Instruction-Level Parallelism


Extract parallelism within a single flow of instructions

- **Overlap the processing of successive instructions**
  - Pipelining
- **Process several instructions at a given time**
  - Overlapping / offload
  - Superscalar processing
- **Apply an instruction on multiple data**
  - SIMD instruction sets

# SIMD Instruction Sets

## Single Instruction Multiple Data

- **General principle**
  - SIMD register == small vector



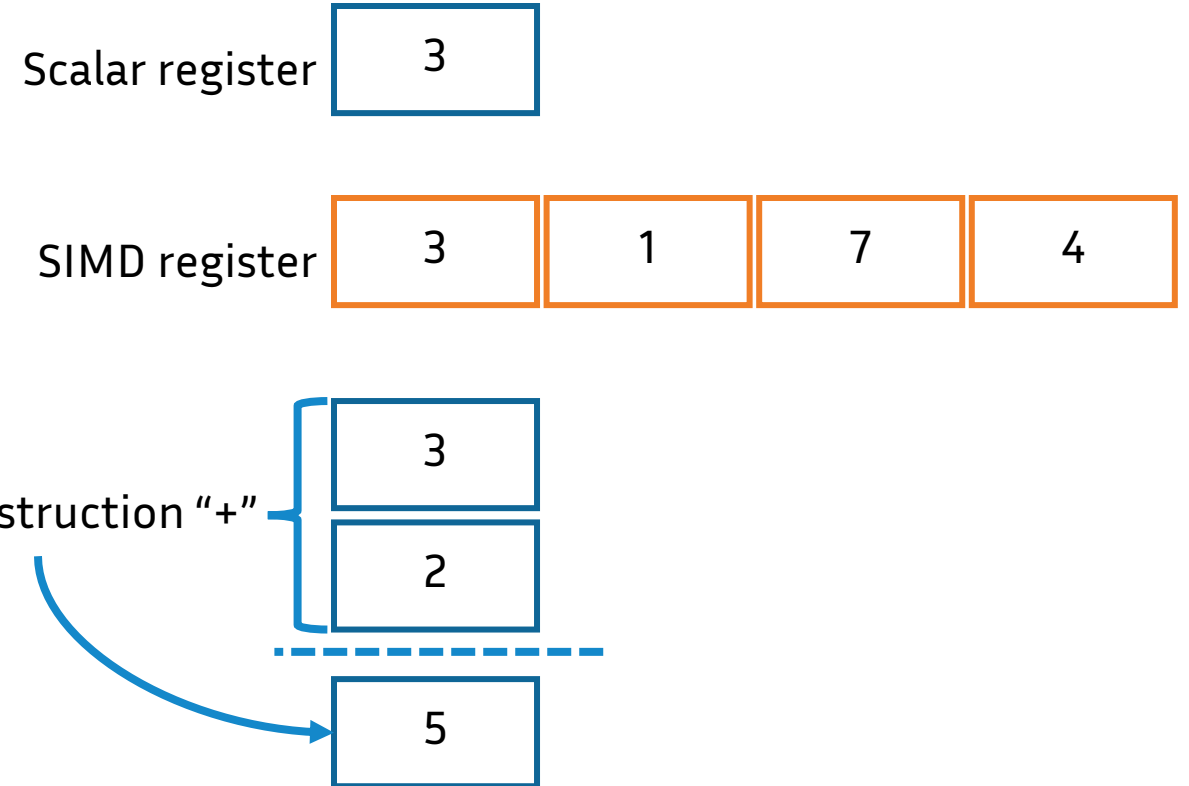
Typical RISC instruction processing sequence

# SIMD Instruction Sets

## Single Instruction Multiple Data

- **General principle**

- SIMD register == small vector
- SIMD arithmetic instructions
  - Ops applied to vector items...
  - ... in parallel



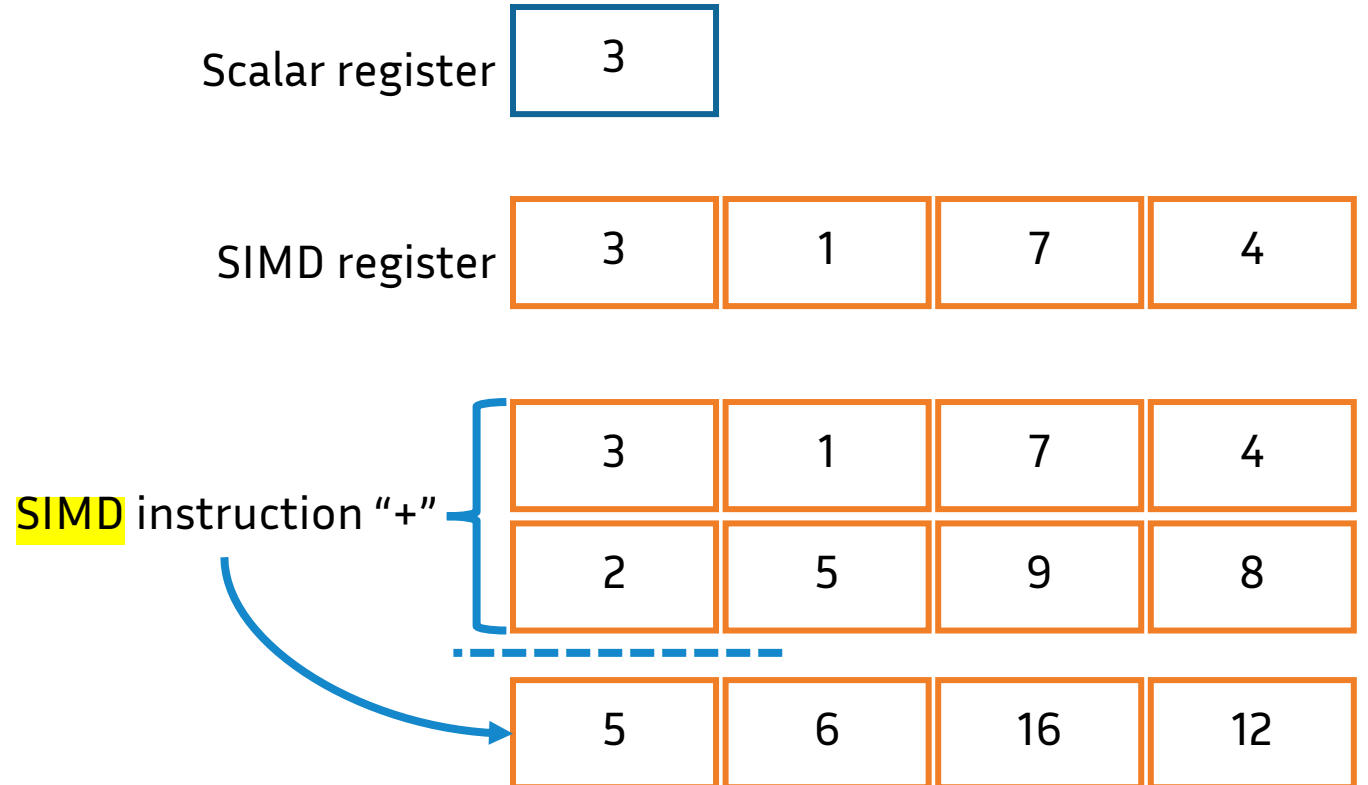
Typical RISC instruction processing sequence

# SIMD Instruction Sets

## Single Instruction Multiple Data

- **General principle**

- SIMD register == small vector
- SIMD arithmetic instructions
  - Ops applied to vector items...
  - ... in parallel



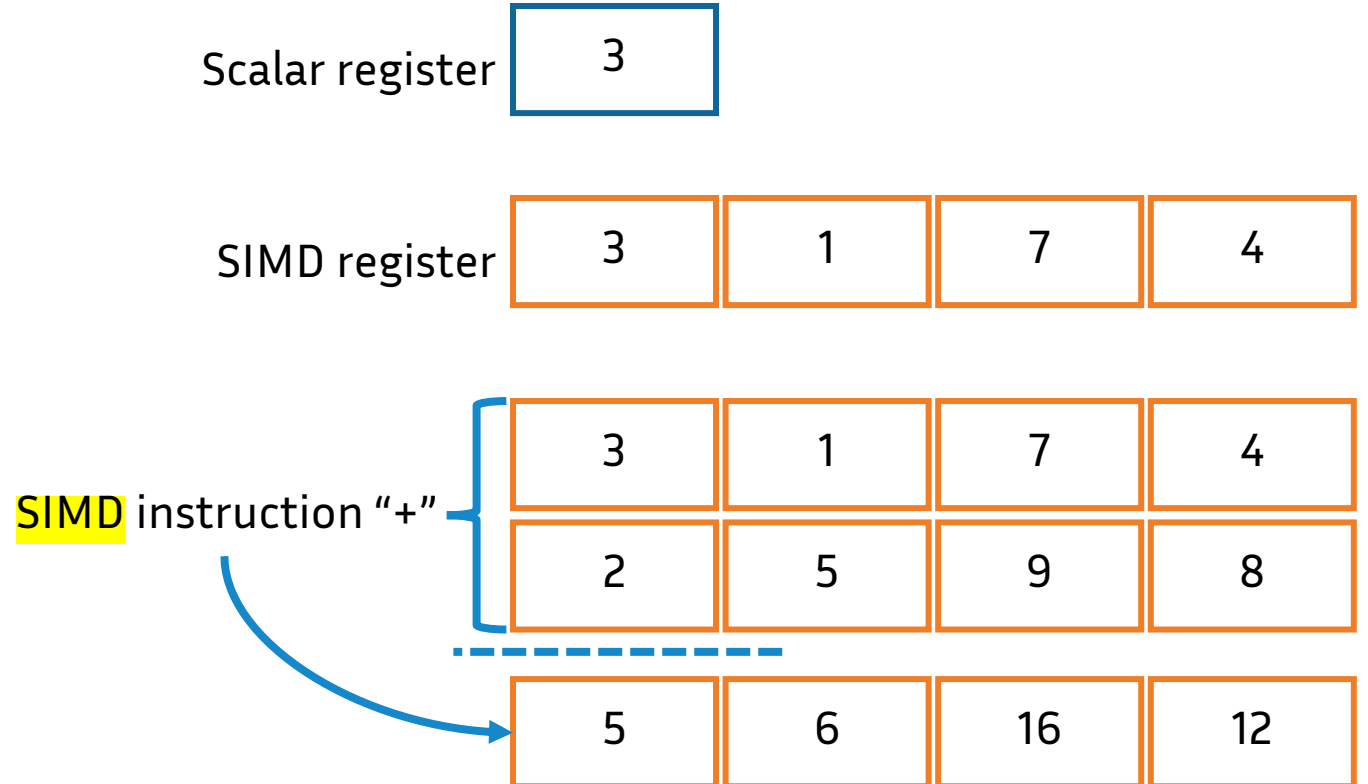
Typical RISC instruction processing sequence



# SIMD Instruction Sets

## Single Instruction Multiple Data

- **General principle**
  - SIMD register == small vector
  - SIMD arithmetic instructions
    - Ops applied to vector items...
    - ... in parallel
- **Example**
  - Intel Intrinsics Guide [\[link\]](#)




Typical RISC instruction processing sequence

# Doing More by Unit of Time

## Hardware parallelism

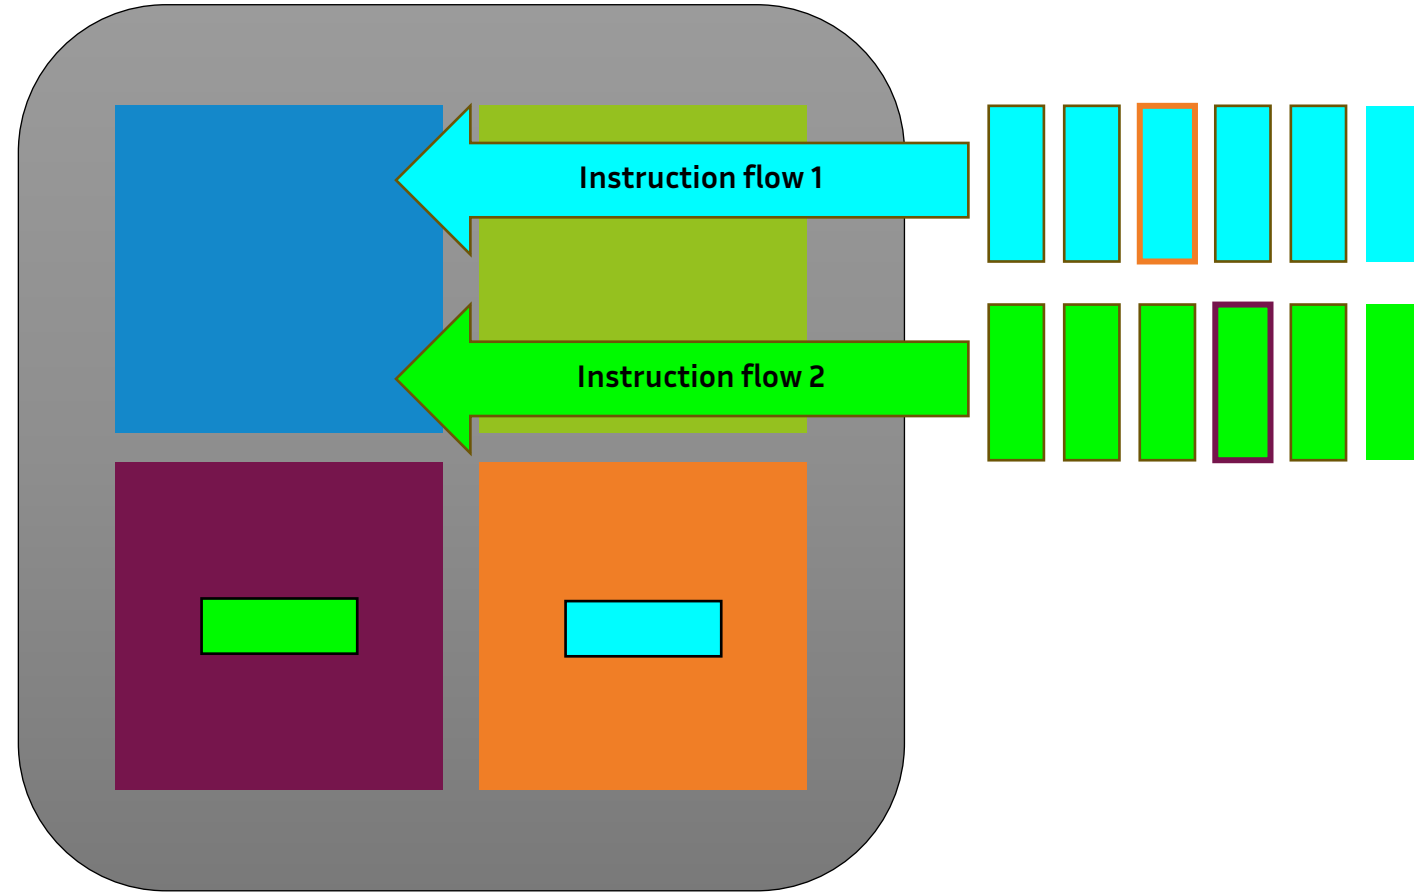
- **Instruction-level parallelism (ILP)**
  - Parallelism within a single instruction flow
- **Thread-level parallelism (TLP)**
  - Parallelism across multiple instruction flows



# Thread-Level Parallelism

Extract parallelism from multiple flows of instructions

- **No automatism**
  - Programs must expose multiple flows




# Parallel Architectures

# Multiprocessor Architectures

Several processors in the same machine

- **Considerations**

- Benefits
- Cost
- Scalability
- Memory




# Multicore Architectures

Several processing areas in the same processor

- **Considerations**

- Benefits
- Cost
- Scalability
- Memory




# Multicore + Multiprocessor Architectures

## Combining the two approaches

- **Considerations**

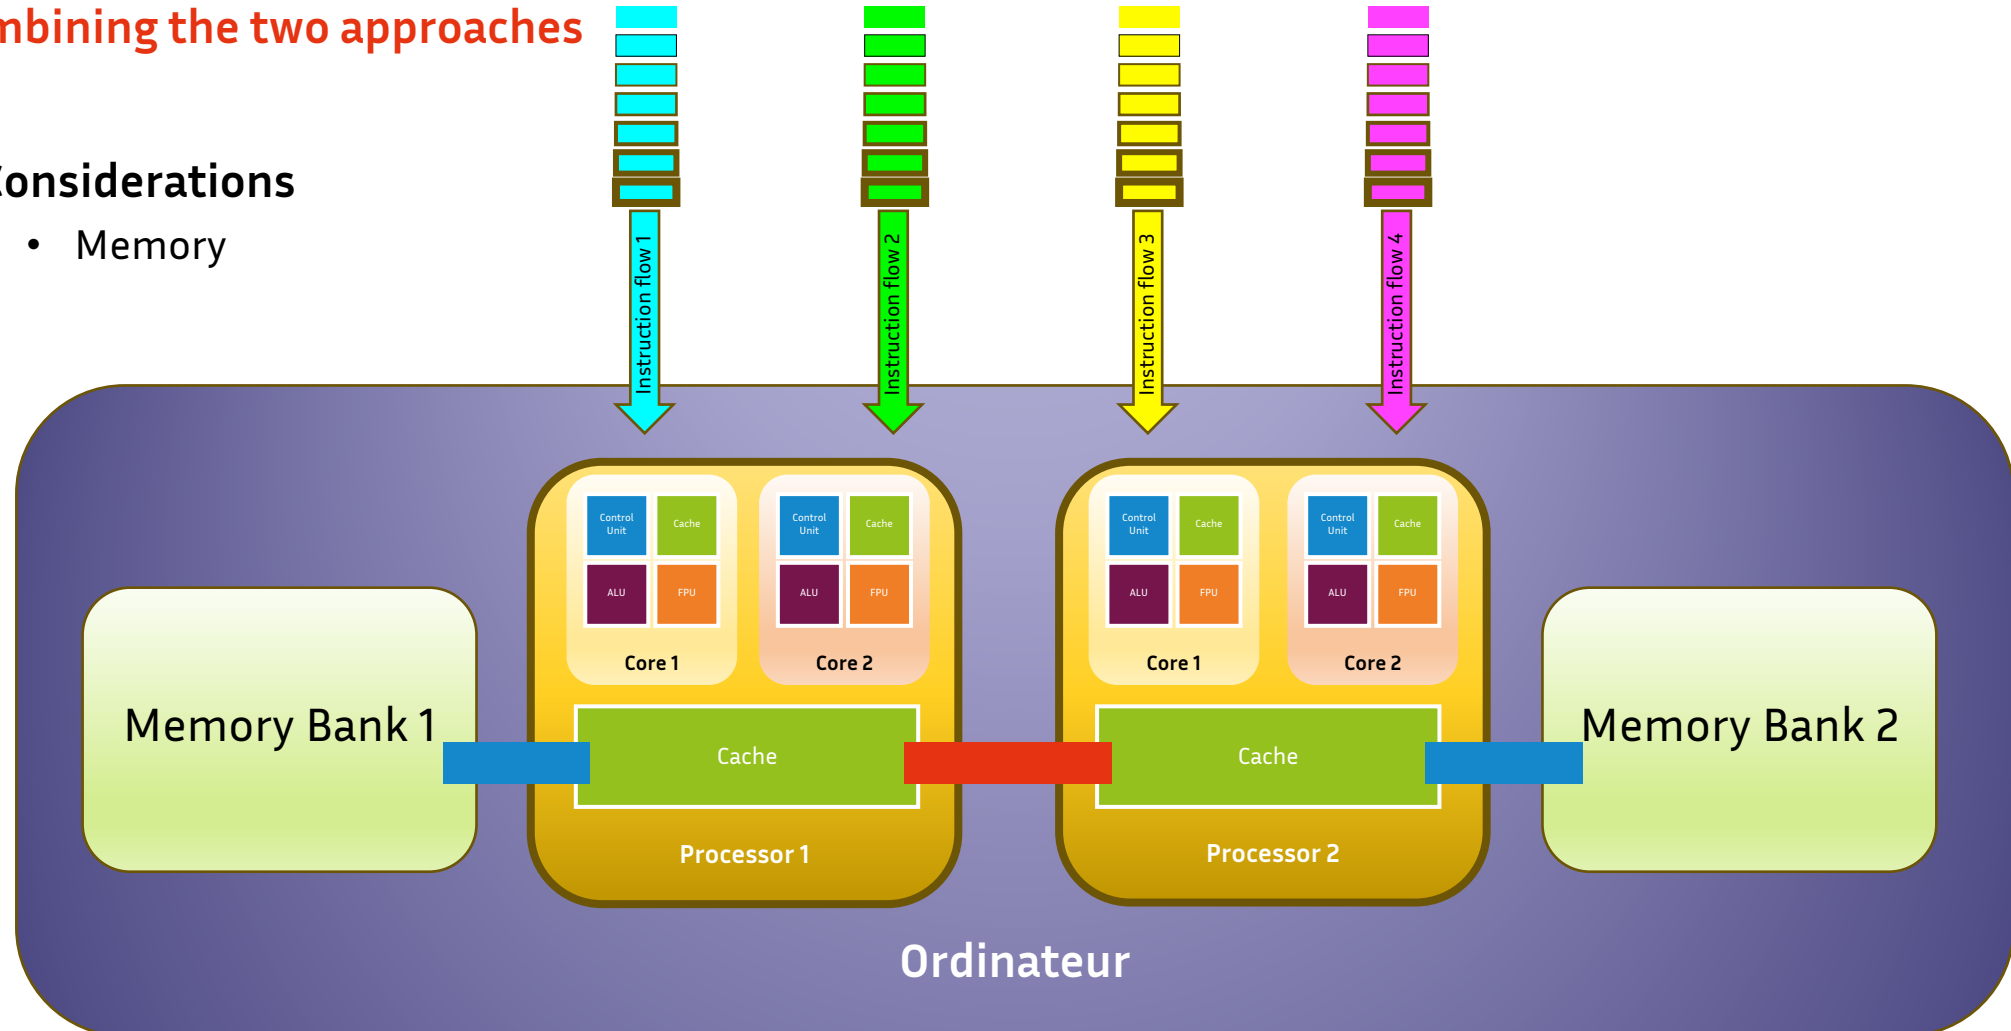
- Benefits
- Cost
- Scalability



# Multicore + Multiprocessor Architectures

## Combining the two approaches

- Considerations
  - Memory



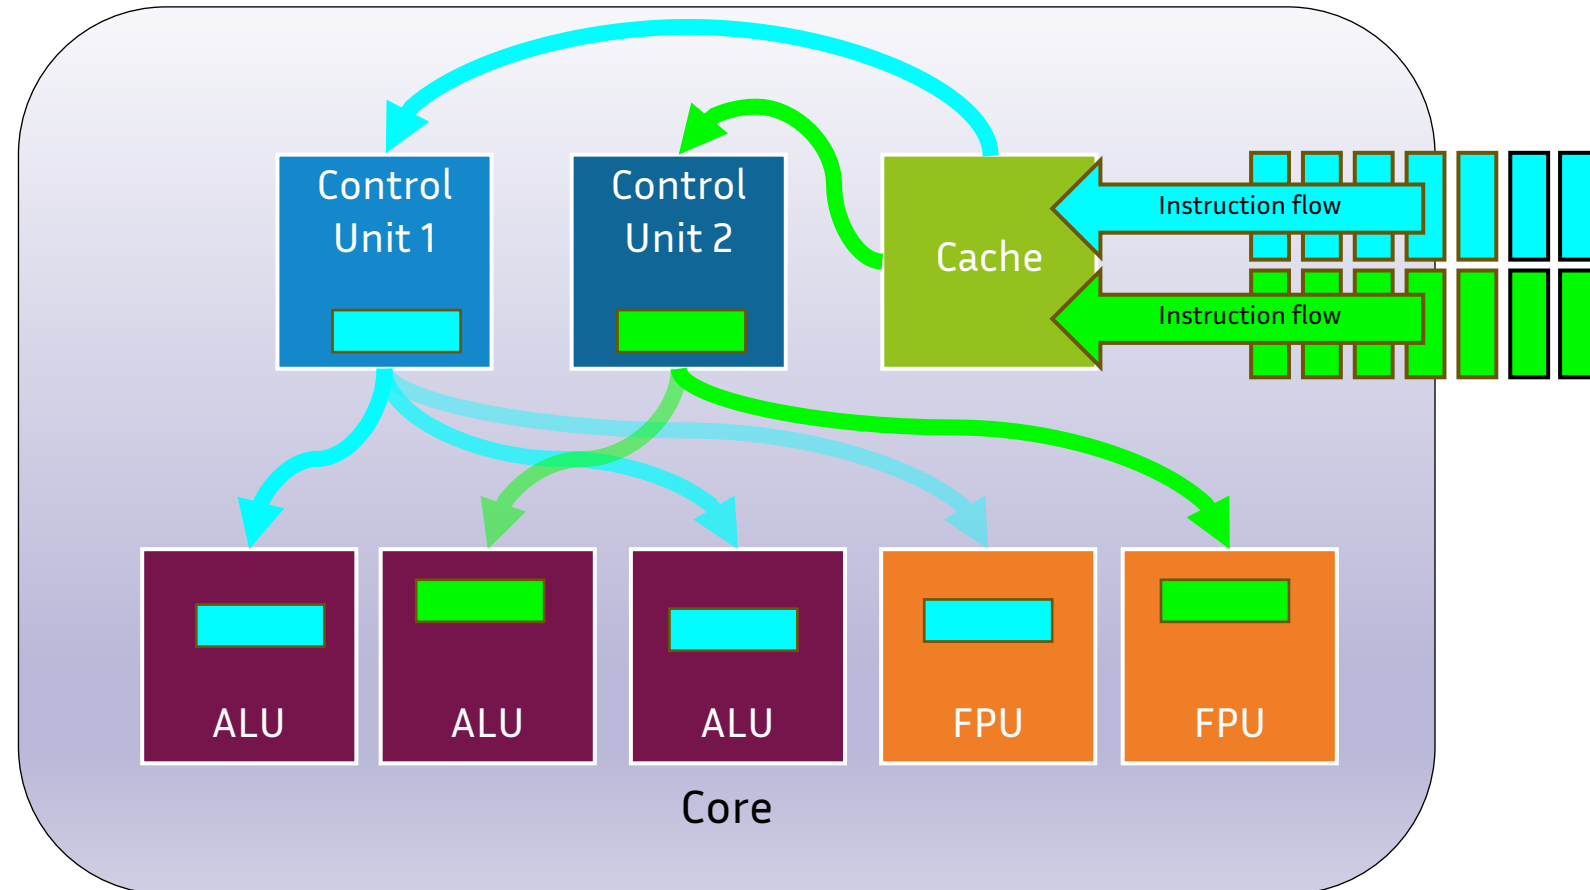


# Hardware Multithreading

## Feeding idle units

- **Principles**

- Multiple thread contexts...
- ... managed in a single core



# Hardware Multithreading


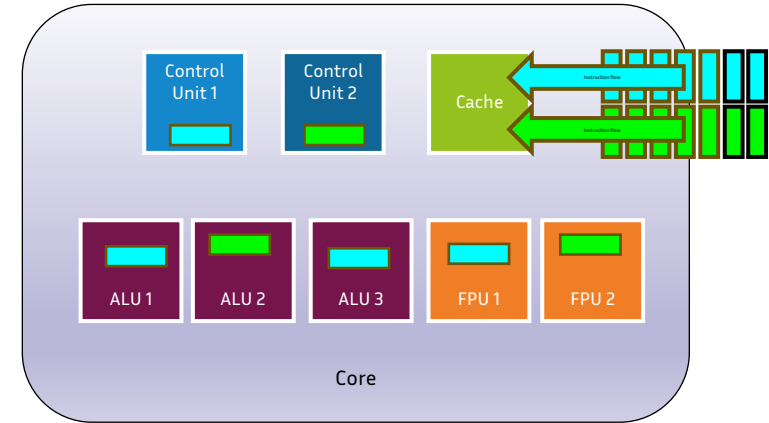
## Feeding idle units

- **Principles**

- Multiple thread contexts...
- ... managed in a single core

- **Considerations**


- Benefits
- Drawbacks



# HW MT + Multicore + Multiprocessor Architectures

Putting everything together

- + SIMD parallelism



# TD

## Topology discovery with libhwloc

- **Assignment, source codes, additional slides**
  - Moodle ENSEIRB, IT390 course
- **Libhwloc on OpenMPI website**
  - [\[link\]](#)
- **Command-line tools**
  - lstopo
  - hwloc-bind
  - hwloc-calc
- **Keywords**
  - Topology
  - Affinity
  - Binding

