

Multicore Architecture Programming – SIMD II

Olivier Aumage

Inria & LaBRI lab.

[olivier.aumage @ inria.fr](mailto:olivier.aumage@inria.fr)

2024 – 2025



Practical Exercises

$$x + y, \alpha.x+y$$

Simple SIMD Arithmetics

- $xpy (x + y)$

- $axpy (\alpha.x+y)$

$$x + y, \alpha.x + y$$

Simple SIMD Arithmetics

- **xpy ($x + y$)**
 - Load, store
 - Addition, element per element
- **axpy ($\alpha.x + y$)**

```
void xpy (const float *x, const float *y, float *z,
          const int vector_size)
{
    /* Check that the vector size is a multiple of the number
       of elements in SIMD registers */
    assert(vector_size % REG_NB_ELEMENTS == 0);

    int i;
    for (i=0; i<vector_size; i+=REG_NB_ELEMENTS)
    {
        __m256 reg_x;
        __m256 reg_y;
        __m256 reg_z;

        /* Load A and B arrays in SIMD registers */
        reg_x = _mm256_load_ps(&x[i]);
        reg_y = _mm256_load_ps(&y[i]);

        /* Perform SIMD add operation */
        reg_z = _mm256_add_ps(reg_x, reg_y);

        /* Store SIMD register in C array */
        _mm256_store_ps(&z[i], reg_z);
    }
}
```

File [xpy.c](#) – routine `xpy()`

$$x + y, \alpha.x + y$$

Simple SIMD Arithmetics

- **xpy ($x + y$)**
 - Load, store
 - Addition, element per element
- **axpy ($\alpha.x + y$)**

```
void xpy (const float *x, const float *y, float *z,
          const int vector_size)
{
    /* Check that the vector size is a multiple of the number
       of elements in SIMD registers */
    assert(vector_size % REG_NB_ELEMENTS == 0);

    int i;
    for (i=0; i<vector_size; i+=REG_NB_ELEMENTS)
    {
        __m256 reg_x;
        __m256 reg_y;
        __m256 reg_z;

        /* Load A and B arrays in SIMD registers */
        reg_x = _mm256_load_ps(&x[i]);
        reg_y = _mm256_load_ps(&y[i]);

        /* Perform SIMD add operation */
        reg_z = _mm256_add_ps(reg_x, reg_y);

        /* Store SIMD register in C array */
        _mm256_store_ps(&z[i], reg_z);
    }
}
```

File [xpy.c](#) – routine `xpy()`

$$x + y, \alpha.x + y$$

Simple SIMD Arithmetics

- **xpy ($x + y$)**
 - Load, store
 - Addition, element per element
- **axpy ($\alpha.x + y$)**

```
void xpy (const float *x, const float *y, float *z,
          const int vector_size)
{
    /* Check that the vector size is a multiple of the number
       of elements in SIMD registers */
    assert(vector_size % REG_NB_ELEMENTS == 0);

    int i;
    for (i=0; i<vector_size; i+=REG_NB_ELEMENTS)
    {
        __m256 reg_x;
        __m256 reg_y;
        __m256 reg_z;

        /* Load A and B arrays in SIMD registers */
        reg_x = _mm256_load_ps(&x[i]);
        reg_y = _mm256_load_ps(&y[i]);

        /* Perform SIMD add operation */
        reg_z = _mm256_add_ps(reg_x, reg_y);

        /* Store SIMD register in C array */
        _mm256_store_ps(&z[i], reg_z);
    }
}
```

File **xpy.c** – routine **xpy()**

$$x + y, \alpha.x + y$$

Simple SIMD Arithmetics

- **xpy ($x + y$)**
 - Load, **store**
 - Addition, element per element
- **axpy ($\alpha.x + y$)**

```
void xpy (const float *x, const float *y, float *z,
          const int vector_size)
{
    /* Check that the vector size is a multiple of the number
       of elements in SIMD registers */
    assert(vector_size % REG_NB_ELEMENTS == 0);

    int i;
    for (i=0; i<vector_size; i+=REG_NB_ELEMENTS)
    {
        __m256 reg_x;
        __m256 reg_y;
        __m256 reg_z;

        /* Load A and B arrays in SIMD registers */
        reg_x = _mm256_load_ps(&x[i]);
        reg_y = _mm256_load_ps(&y[i]);

        /* Perform SIMD add operation */
        reg_z = _mm256_add_ps(reg_x, reg_y);

        /* Store SIMD register in C array */
        _mm256_store_ps(&z[i], reg_z);
    }
}
```

File **xpy.c** – routine **xpy()**

$$x + y, \alpha.x + y$$

Simple SIMD Arithmetics

- **xpy ($x + y$)**
 - Load, store
 - Addition, element per element
- **axpy ($\alpha.x + y$)**

```
void xpy (const float *x, const float *y, float *z,
          const int vector_size)
{
    /* Check that the vector size is a multiple of the number
       of elements in SIMD registers */
    assert(vector_size % REG_NB_ELEMENTS == 0);

    int i;
    for (i=0; i<vector_size; i+=REG_NB_ELEMENTS)
    {
        __m256 reg_x;
        __m256 reg_y;
        __m256 reg_z;

        /* Load A and B arrays in SIMD registers */
        reg_x = _mm256_load_ps(&x[i]);
        reg_y = _mm256_load_ps(&y[i]);

        /* Perform SIMD add operation */
        reg_z = _mm256_add_ps(reg_x, reg_y);

        /* Store SIMD register in C array */
        _mm256_store_ps(&z[i], reg_z);
    }
}
```

File [xpy.c](#) – routine `xpy()`

$$x + y, \alpha.x + y$$

Simple SIMD Arithmetics

- **xpy ($x + y$)**
 - Load, store
 - Addition, element per element
- **axpy ($\alpha.x + y$)**
 - Initialization with a constant (set1)
 - Fused Multiply-Accumulate (FMA)

```
void axpy (const float alpha, const float *x, const float *y,
          float *z, const int vector_size)
{
    /* Check that the vector size is a multiple of the number of
     * elements in SIMD registers */
    assert(vector_size % REG_NB_ELEMENTS == 0);

    __m256 reg_a = _mm256_set1_ps(alpha);

    int i;
    for (i=0; i<vector_size; i+=REG_NB_ELEMENTS)
    {
        __m256 reg_x;
        __m256 reg_y;
        __m256 reg_z;

        /* Load A and B arrays in SIMD registers */
        reg_x = _mm256_load_ps(&x[i]);
        reg_y = _mm256_load_ps(&y[i]);

        /* Perform SIMD add operation */
        reg_z = _mm256_fmadd_ps(reg_a, reg_x, reg_y);

        /* Store SIMD register in C array */
        _mm256_store_ps(&z[i], reg_z);
    }
}
```

File [axpy.c](#) – routine `axpy()`

$$x + y, \alpha.x + y$$

Simple SIMD Arithmetics

- **xpy ($x + y$)**
 - Load, store
 - Addition, element per element
- **axpy ($\alpha.x + y$)**
 - Initialization with a constant (set1)
 - Fused Multiply-Accumulate (FMA)

```
void axpy (const float alpha, const float *x, const float *y,
           float *z, const int vector_size)
{
    /* Check that the vector size is a multiple of the number of
     * elements in SIMD registers */
    assert(vector_size % REG_NB_ELEMENTS == 0);

    __m256 reg_a = _mm256_set1_ps(alpha);

    int i;
    for (i=0; i<vector_size; i+=REG_NB_ELEMENTS)
    {
        __m256 reg_x;
        __m256 reg_y;
        __m256 reg_z;

        /* Load A and B arrays in SIMD registers */
        reg_x = _mm256_load_ps(&x[i]);
        reg_y = _mm256_load_ps(&y[i]);

        /* Perform SIMD add operation */
        reg_z = _mm256_fmadd_ps(reg_a, reg_x, reg_y);

        /* Store SIMD register in C array */
        _mm256_store_ps(&z[i], reg_z);
    }
}
```

File [axpy.c](#) – routine `axpy()`

$$x + y, \alpha.x + y$$

Simple SIMD Arithmetics

- **xpy ($x + y$)**
 - Load, store
 - Addition, element per element
- **axpy ($\alpha.x + y$)**
 - Initialization with a constant (set1)
 - Fused Multiply-Accumulate (FMA)

```
void axpy (const float alpha, const float *x, const float *y,
           float *z, const int vector_size)
{
    /* Check that the vector size is a multiple of the number of
     * elements in SIMD registers */
    assert(vector_size % REG_NB_ELEMENTS == 0);

    __m256 reg_a = _mm256_set1_ps(alpha);

    int i;
    for (i=0; i<vector_size; i+=REG_NB_ELEMENTS)
    {
        __m256 reg_x;
        __m256 reg_y;
        __m256 reg_z;

        /* Load A and B arrays in SIMD registers */
        reg_x = _mm256_load_ps(&x[i]);
        reg_y = _mm256_load_ps(&y[i]);

        /* Perform SIMD add operation */
        reg_z = _mm256_fmadd_ps(reg_a, reg_x, reg_y);

        /* Store SIMD register in C array */
        _mm256_store_ps(&z[i], reg_z);
    }
}
```

File [axpy.c](#) – routine `axpy()`

Reduction, Dot Product

Slightly Less Simple SIMD Arithmetics

- Reduction

Reduction, Dot Product

Slightly Less Simple SIMD Arithmetics

- **Reduction**

- Permute

```
void reduce_sum (const float *x, const int vector_size, __m256 *r)
{
    assert(vector_size % REG_NB_ELEMENTS == 0);

    __m256 a;
    if (vector_size > 0)
    {
        a = _mm256_load_ps(&x[0]);

        int i;
        for (i=REG_NB_ELEMENTS; i<vector_size; i+=REG_NB_ELEMENTS)
        {
            __m256 v = _mm256_load_ps(&x[i]);
            a = _mm256_add_ps(a, v);
        }

        /* Add elements in the lower 128-bit and the elts in the higher 128-bit
         * This sums element i with element i + 4 */
        a = _mm256_add_ps(a, _mm256_permute2f128_ps(a, a, _MM_SHUFFLE(0,0,0,1)));

        /* This sums element i with element i + 2 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(1,0,3,2)));

        /* This sums element i with element i + 1 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(2,3,0,1)));
    }
    else
    {
        a = _mm256_setzero_ps();
    }
    *r = a;
}
```

File `reduce_sum.c` – routine `reduce_sum()`

Reduction, Dot Product

Slightly Less Simple SIMD Arithmetics

- **Reduction**

- Permute

```
void reduce_sum (const float *x, const int vector_size, __m256 *r)
{
    assert(vector_size % REG_NB_ELEMENTS == 0);

    __m256 a;
    if (vector_size > 0)
    {
        a = _mm256_load_ps(&x[0]);

        int i;
        for (i=REG_NB_ELEMENTS; i<vector_size; i+=REG_NB_ELEMENTS)
        {
            __m256 v = _mm256_load_ps(&x[i]);
            a = _mm256_add_ps(a, v);
        }

        /* Add elements in the lower 128-bit and the elts in the higher 128-bit
         * This sums element i with element i + 4 */
        a = _mm256_add_ps(a, _mm256_permute2f128_ps(a, a, _MM_SHUFFLE(0,0,0,1)));

        /* This sums element i with element i + 2 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(1,0,3,2)));

        /* This sums element i with element i + 1 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(2,3,0,1)));
    }
    else
    {
        a = _mm256_setzero_ps();
    }
    *r = a;
}
```

File `reduce_sum.c` – routine `reduce_sum()`

Reduction, Dot Product

Slightly Less Simple SIMD Arithmetics

- **Reduction**

- Permute

```
void reduce_sum (const float *x, const int vector_size, __m256 *r)
{
    assert(vector_size % REG_NB_ELEMENTS == 0);

    __m256 a;
    if (vector_size > 0)
    {
        a = _mm256_load_ps(&x[0]);

        int i;
        for (i=REG_NB_ELEMENTS; i<vector_size; i+=REG_NB_ELEMENTS)
        {
            __m256 v = _mm256_load_ps(&x[i]);
            a = _mm256_add_ps(a, v);
        }

        /* Add elements in the lower 128-bit and the elts in the higher 128-bit
         * This sums element i with element i + 4 */
        a = _mm256_add_ps(a, _mm256_permute2f128_ps(a, a, _MM_SHUFFLE(0,0,0,1)));

        /* This sums element i with element i + 2 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(1,0,3,2)));

        /* This sums element i with element i + 1 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(2,3,0,1)));
    }
    else
    {
        a = _mm256_setzero_ps();
    }
    *r = a;
}
```

File `reduce_sum.c` – routine `reduce_sum()`

Reduction, Dot Product

Slightly Less Simple SIMD Arithmetics

- **Reduction**

- Permute



```
void reduce_sum (const float *x, const int vector_size, __m256 *r)
{
    assert(vector_size % REG_NB_ELEMENTS == 0);

    __m256 a;
    if (vector_size > 0)
    {
        a = _mm256_load_ps(&x[0]);

        int i;
        for (i=REG_NB_ELEMENTS; i<vector_size; i+=REG_NB_ELEMENTS)
        {
            __m256 v = _mm256_load_ps(&x[i]);
            a = _mm256_add_ps(a, v);
        }

        /* Add elements in the lower 128-bit and the elts in the higher 128-bit
         * This sums element i with element i + 4 */
        a = _mm256_add_ps(a, _mm256_permute2f128_ps(a, a, _MM_SHUFFLE(0,0,0,1)));

        /* This sums element i with element i + 2 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(1,0,3,2)));

        /* This sums element i with element i + 1 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(2,3,0,1)));
    }
    else
    {
        a = _mm256_setzero_ps();
    }
    *r = a;
}
```

File `reduce_sum.c` – routine `reduce_sum()`

Reduction, Dot Product

Slightly Less Simple SIMD Arithmetics

- **Reduction**

- Permute

Elements (32-bit, single precision floating point)



```
void reduce_sum (const float *x, const int vector_size, __m256 *r)
{
    assert(vector_size % REG_NB_ELEMENTS == 0);

    __m256 a;
    if (vector_size > 0)
    {
        a = _mm256_load_ps(&x[0]);

        int i;
        for (i=REG_NB_ELEMENTS; i<vector_size; i+=REG_NB_ELEMENTS)
        {
            __m256 v = _mm256_load_ps(&x[i]);
            a = _mm256_add_ps(a, v);
        }

        /* Add elements in the lower 128-bit and the elts in the higher 128-bit
         * This sums element i with element i + 4 */
        a = _mm256_add_ps(a, _mm256_permute2f128_ps(a, a, _MM_SHUFFLE(0,0,0,1)));

        /* This sums element i with element i + 2 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(1,0,3,2)));

        /* This sums element i with element i + 1 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(2,3,0,1)));
    }
    else
    {
        a = _mm256_setzero_ps();
    }
    *r = a;
}
```

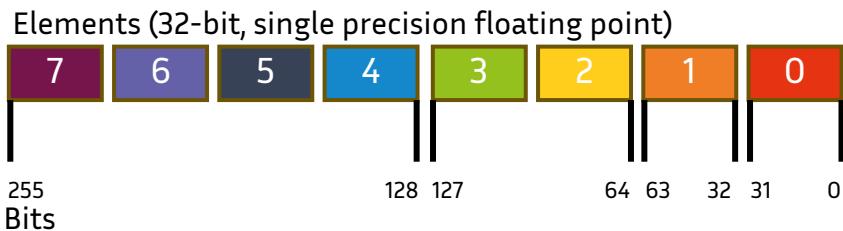
File `reduce_sum.c` – routine `reduce_sum()`

Reduction, Dot Product

Slightly Less Simple SIMD Arithmetics

- **Reduction**

- Permute



```
void reduce_sum (const float *x, const int vector_size, __m256 *r)
{
    assert(vector_size % REG_NB_ELEMENTS == 0);

    __m256 a;
    if (vector_size > 0)
    {
        a = _mm256_load_ps(&x[0]);

        int i;
        for (i=REG_NB_ELEMENTS; i<vector_size; i+=REG_NB_ELEMENTS)
        {
            __m256 v = _mm256_load_ps(&x[i]);
            a = _mm256_add_ps(a, v);
        }

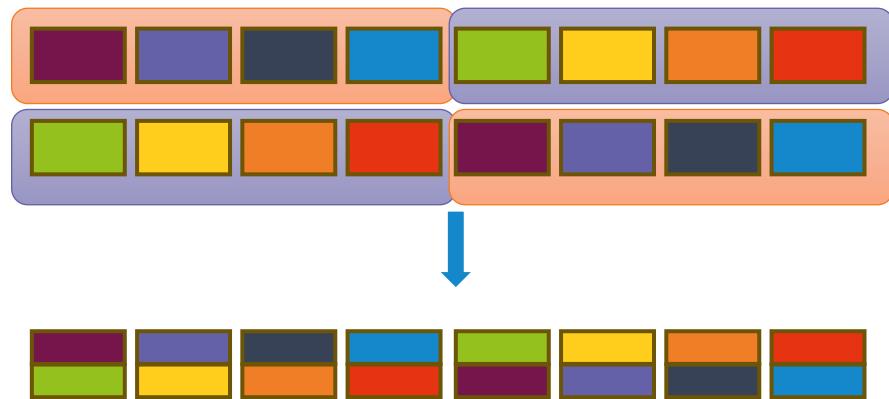
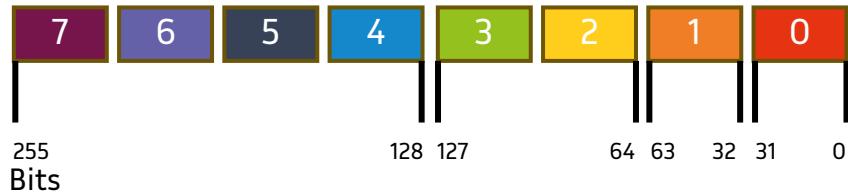
        /* Add elements in the lower 128-bit and the elts in the higher 128-bit
         * This sums element i with element i + 4 */
        a = _mm256_add_ps(a, _mm256_permute2f128_ps(a, a, _MM_SHUFFLE(0,0,0,1)));

        /* This sums element i with element i + 2 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(1,0,3,2)));

        /* This sums element i with element i + 1 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(2,3,0,1)));
    }
    else
    {
        a = _mm256_setzero_ps();
    }
    *r = a;
}
```

File `reduce_sum.c` – routine `reduce_sum()`

Reduction, Dot Product



```
void reduce_sum (const float *x, const int vector_size, __m256 *r)
{
    assert(vector_size % REG_NB_ELEMENTS == 0);

    __m256 a;
    if (vector_size > 0)
    {
        a = _mm256_load_ps(&x[0]);

        int i;
        for (i=REG_NB_ELEMENTS; i<vector_size; i+=REG_NB_ELEMENTS)
        {
            __m256 v = _mm256_load_ps(&x[i]);
            a = _mm256_add_ps(a, v);
        }

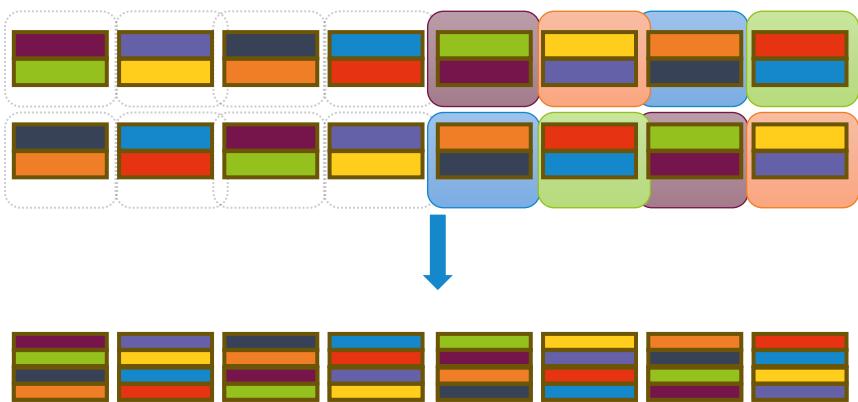
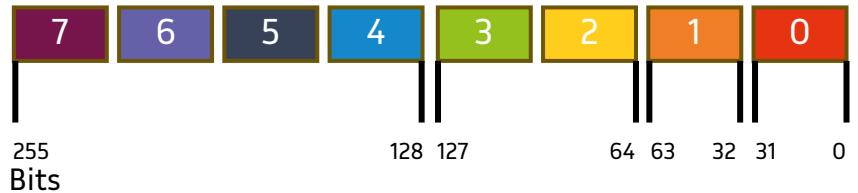
        /* Add elements in the lower 128-bit and the elts in the higher 128-bit
         * This sums element i with element i + 4 */
        a = _mm256_add_ps(a, _mm256_permute2f128_ps(a, a, _MM_SHUFFLE(0,0,0,1)));

        /* This sums element i with element i + 2 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(1,0,3,2)));

        /* This sums element i with element i + 1 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(2,3,0,1)));
    }
    else
    {
        a = _mm256_setzero_ps();
    }
    *r = a;
}
```

File `reduce_sum.c` – routine `reduce_sum()`

Reduction, Dot Product



```
void reduce_sum (const float *x, const int vector_size, __m256 *r)
{
    assert(vector_size % REG_NB_ELEMENTS == 0);

    __m256 a;
    if (vector_size > 0)
    {
        a = _mm256_load_ps(&x[0]);

        int i;
        for (i=REG_NB_ELEMENTS; i<vector_size; i+=REG_NB_ELEMENTS)
        {
            __m256 v = _mm256_load_ps(&x[i]);
            a = _mm256_add_ps(a, v);
        }

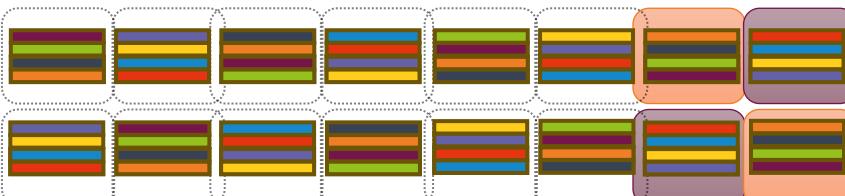
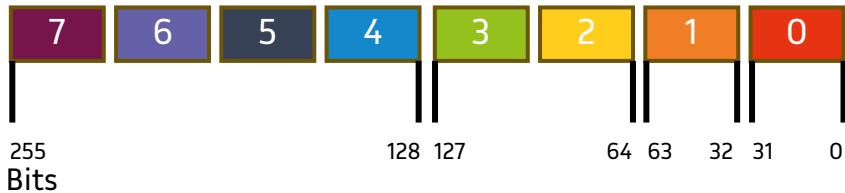
        /* Add elements in the lower 128-bit and the elts in the higher 128-bit
         * This sums element i with element i + 4 */
        a = _mm256_add_ps(a, _mm256_permute2f128_ps(a, a, _MM_SHUFFLE(0,0,0,1)));

        /* This sums element i with element i + 2 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(1,0,3,2)));

        /* This sums element i with element i + 1 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(2,3,0,1)));
    }
    else
    {
        a = _mm256_setzero_ps();
    }
    *r = a;
}
```

File `reduce_sum.c` – routine `reduce_sum()`

Reduction, Dot Product



```
void reduce_sum (const float *x, const int vector_size, __m256 *r)
{
    assert(vector_size % REG_NB_ELEMENTS == 0);

    __m256 a;
    if (vector_size > 0)
    {
        a = _mm256_load_ps(&x[0]);

        int i;
        for (i=REG_NB_ELEMENTS; i<vector_size; i+=REG_NB_ELEMENTS)
        {
            __m256 v = _mm256_load_ps(&x[i]);
            a = _mm256_add_ps(a, v);
        }

        /* Add elements in the lower 128-bit and the elts in the higher 128-bit
         * This sums element i with element i + 4 */
        a = _mm256_add_ps(a, _mm256_permute2f128_ps(a, a, _MM_SHUFFLE(0,0,0,1)));

        /* This sums element i with element i + 2 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(1,0,3,2)));

        /* This sums element i with element i + 1 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(2,3,0,1)));
    }
    else
    {
        a = _mm256_setzero_ps();
    }
    *r = a;
}
```

File `reduce_sum.c` – routine `reduce_sum()`

Reduction, Dot Product

Slightly Less Simple SIMD Arithmetics

- **Reduction**

- Permute

- **Dot product**

```
void reduce_sum (const float *x, const int vector_size, __m256 *r)
{
    assert(vector_size % REG_NB_ELEMENTS == 0);

    __m256 a;
    if (vector_size > 0)
    {
        a = _mm256_load_ps(&x[0]);

        int i;
        for (i=REG_NB_ELEMENTS; i<vector_size; i+=REG_NB_ELEMENTS)
        {
            __m256 v = _mm256_load_ps(&x[i]);
            a = _mm256_add_ps(a, v);
        }

        /* Add elements in the lower 128-bit and the elts in the higher 128-bit
         * This sums element i with element i + 4 */
        a = _mm256_add_ps(a, _mm256_permute2f128_ps(a, a, _MM_SHUFFLE(0,0,0,1)));

        /* This sums element i with element i + 2 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(1,0,3,2)));

        /* This sums element i with element i + 1 */
        a = _mm256_add_ps(a, _mm256_permute_ps(a, _MM_SHUFFLE(2,3,0,1)));
    }
    else
    {
        a = _mm256_setzero_ps();
    }
    *r = a;
}
```

File `reduce_sum.c` – routine `reduce_sum()`

Reduction, Dot Product

Slightly Less Simple SIMD A

- Reduction
 - Permute
- Dot product
 - FMA
 - Reduction

```
float dot_product (const float *x, const float *y, const int vector_size)
{
    assert(vector_size % REG_NB_ELEMENTS == 0);
    float result = 0;

    if (vector_size > 0)
    {
        __m256 reg_acc = _mm256_setzero_ps();

        int i;
        for (i=0; i<vector_size; i+=REG_NB_ELEMENTS)
        {
            __m256 reg_x;
            __m256 reg_y;

            /* Load A and B arrays in SIMD registers */
            reg_x = _mm256_load_ps(&x[i]);
            reg_y = _mm256_load_ps(&y[i]);

            /* Perform SIMD add operation */
            reg_acc = _mm256_fmadd_ps(reg_x, reg_y, reg_acc);
        }

        /* Perform the sum reduction */
        reg_acc = _mm256_add_ps(reg_acc, _mm256_permute2f128_ps(reg_acc, reg_acc, _MM_SHUFFLE(0,0,0,1)));
        reg_acc = _mm256_add_ps(reg_acc, _mm256_permute_ps(reg_acc, _MM_SHUFFLE(1,0,3,2)));
        reg_acc = _mm256_add_ps(reg_acc, _mm256_permute_ps(reg_acc, _MM_SHUFFLE(2,3,0,1)));

        result = reg_acc[0];
    }
    return result;
}
```

File `dotp.c` – routine `dot_product()`

Reduction, Dot Product

Slightly Less Simple SIMD A

- Reduction
 - Permute
- Dot product
 - FMA
 - Reduction

```
float dot_product (const float *x, const float *y, const int vector_size)
{
    assert(vector_size % REG_NB_ELEMENTS == 0);
    float result = 0;

    if (vector_size > 0)
    {
        __m256 reg_acc = _mm256_setzero_ps();

        int i;
        for (i=0; i<vector_size; i+=REG_NB_ELEMENTS)
        {
            __m256 reg_x;
            __m256 reg_y;

            /* Load A and B arrays in SIMD registers */
            reg_x = _mm256_load_ps(&x[i]);
            reg_y = _mm256_load_ps(&y[i]);

            /* Perform SIMD add operation */
            reg_acc = _mm256_fmadd_ps(reg_x, reg_y, reg_acc);
        }

        /* Perform the sum reduction */
        reg_acc = _mm256_add_ps(reg_acc, _mm256_permute2f128_ps(reg_acc, reg_acc, _MM_SHUFFLE(0,0,0,1)));
        reg_acc = _mm256_add_ps(reg_acc, _mm256_permute_ps(reg_acc, _MM_SHUFFLE(1,0,3,2)));
        reg_acc = _mm256_add_ps(reg_acc, _mm256_permute_ps(reg_acc, _MM_SHUFFLE(2,3,0,1)));

        result = reg_acc[0];
    }
    return result;
}
```

File `dotp.c` – routine `dot_product()`

Reduction, Dot Product

Slightly Less Simple SIMD A

- Reduction
 - Permute
- Dot product
 - FMA
 - Reduction

```
float dot_product (const float *x, const float *y, const int vector_size)
{
    assert(vector_size % REG_NB_ELEMENTS == 0);
    float result = 0;

    if (vector_size > 0)
    {
        __m256 reg_acc = _mm256_setzero_ps();

        int i;
        for (i=0; i<vector_size; i+=REG_NB_ELEMENTS)
        {
            __m256 reg_x;
            __m256 reg_y;

            /* Load A and B arrays in SIMD registers */
            reg_x = _mm256_load_ps(&x[i]);
            reg_y = _mm256_load_ps(&y[i]);

            /* Perform SIMD add operation */
            reg_acc = _mm256_fmadd_ps(reg_x, reg_y, reg_acc);
        }

        /* Perform the sum reduction */
        reg_acc = _mm256_add_ps(reg_acc, _mm256_permute2f128_ps(reg_acc, reg_acc, _MM_SHUFFLE(0,0,0,1)));
        reg_acc = _mm256_add_ps(reg_acc, _mm256_permute_ps(reg_acc, _MM_SHUFFLE(1,0,3,2)));
        reg_acc = _mm256_add_ps(reg_acc, _mm256_permute_ps(reg_acc, _MM_SHUFFLE(2,3,0,1)));

        result = reg_acc[0];
    }
    return result;
}
```

File `dotp.c` – routine `dot_product()`

Reduction, Dot Product

Slightly Less Simple SIMD A

- Reduction
 - Permute
- Dot product
 - FMA
 - Reduction

```
float dot_product (const float *x, const float *y, const int vector_size)
{
    assert(vector_size % REG_NB_ELEMENTS == 0);
    float result = 0;

    if (vector_size > 0)
    {
        __m256 reg_acc = _mm256_setzero_ps();

        int i;
        for (i=0; i<vector_size; i+=REG_NB_ELEMENTS)
        {
            __m256 reg_x;
            __m256 reg_y;

            /* Load A and B arrays in SIMD registers */
            reg_x = _mm256_load_ps(&x[i]);
            reg_y = _mm256_load_ps(&y[i]);

            /* Perform SIMD add operation */
            reg_acc = _mm256_fmadd_ps(reg_x, reg_y, reg_acc);
        }

        /* Perform the sum reduction */
        reg_acc = _mm256_add_ps(reg_acc, _mm256_permute2f128_ps(reg_acc, reg_acc, _MM_SHUFFLE(0,0,0,1)));
        reg_acc = _mm256_add_ps(reg_acc, _mm256_permute_ps(reg_acc, _MM_SHUFFLE(1,0,3,2)));
        reg_acc = _mm256_add_ps(reg_acc, _mm256_permute_ps(reg_acc, _MM_SHUFFLE(2,3,0,1)));

        result = reg_acc[0];
    }
    return result;
}
```

File `dotp.c` – routine `dot_product()`

Minmax

Comparisons and Masks

- minmax version 1

Minmax

~~Comparisons and Masks~~

- **minmax version 1**

- Dedicated routines
- Always check whether a function already exists! 😊

```
void minmax(__m256i x, __m256i y, __m256i *zmin, __m256i *zmax)
{
    *zmax = _mm256_max_epi32(x, y);
    *zmin = _mm256_min_epi32(x, y);
}
```

File [minmax.c](#) – routine `minmax()`

Minmax

Comparisons and Masks

- **minmax version 2**
 - *For illustration purpose*
 - Comparisons
 - Masks & Boolean operations

```
void minmax2(__m256i x, __m256i y, __m256i *zmin, __m256i *zmax)
{
    __m256i c = _mm256_cmpgt_epi32(x, y);

    *zmax = _mm256_or_si256(_mm256_and_si256(c,x), _mm256_andnot_si256(c,y));
    *zmin = _mm256_or_si256(_mm256_andnot_si256(c,x), _mm256_and_si256(c,y));
}
```

File [minmax.c](#) – routine `minmax()`

Minmax

Comparisons and Masks

- **minmax version 2**

- *For illustration purpose*
- Comparisons
- Masks & Boolean operations

```
void minmax2(__m256i x, __m256i y, __m256i *zmin, __m256i *zmax)
{
    __m256i c = _mm256_cmpgt_epi32(x, y);

    *zmax = _mm256_or_si256(
        _mm256_and_si256(c, x),
        _mm256_andnot_si256(c, y));

    *zmin = _mm256_or_si256(
        _mm256_andnot_si256(c, x),
        _mm256_and_si256(c, y));
}
```

File `minmax.c` — routine `minmax()`

x	-3	-1	4	-12	7	-7	15	2
y	-7	18	3	4	-1	-19	11	-10

Minmax

Comparisons and Masks

- minmax version 2
 - For illustration purpose
 - Comparisons
 - Masks & Boolean operations

```
void minmax2(__m256i x, __m256i y, __m256i *zmin, __m256i *zmax)
{
    __m256i c = _mm256_cmpgt_epi32(x, y);

    *zmax = _mm256_or_si256(
        _mm256_and_si256(c, x),
        _mm256_andnot_si256(c, y));

    *zmin = _mm256_or_si256(
        _mm256_andnot_si256(c, x),
        _mm256_and_si256(c, y));
}
```

File minmax.c – routine minmax()

x	-3	-1	4	-12	7	-7	15	2
y	-7	18	3	4	-1	-19	11	-10
x > y	c	0xFFFFFFFF	0x00000000	0xFFFFFFFF	0x00000000	0xFFFFFFFF	0xFFFFFFFF	0xFFFFFFFF

Minmax

Comparisons and Masks

- minmax version 2

- For illustration purpose
- Comparisons
- Masks & Boolean operations

```
void minmax2(__m256i x, __m256i y, __m256i *zmin, __m256i *zmax)
{
    __m256i c = _mm256_cmpgt_epi32(x, y);

    *zmax = _mm256_or_si256(
        _mm256_and_si256(c, x),
        _mm256_andnot_si256(c, y));

    *zmin = _mm256_or_si256(
        _mm256_andnot_si256(c, x),
        _mm256_and_si256(c, y));
}
```

File minmax.c – routine minmax()

x	-3	-1	4	-12	7	-7	15	2	
c	0xFFFFFFFF	0x00000000	0xFFFFFFFF	0x00000000	0xFFFFFFFF	0xFFFFFFFF	0xFFFFFFFF	0xFFFFFFFF	
c and x	Result 1	-3	0	4	0	7	-7	15	2

Minmax

Comparisons and Masks

- minmax version 2

- For illustration purpose
- Comparisons
- Masks & Boolean operations

```
void minmax2(__m256i x, __m256i y, __m256i *zmin, __m256i *zmax)
{
    __m256i c = _mm256_cmpgt_epi32(x, y);

    *zmax = _mm256_or_si256(
        _mm256_and_si256(c, x),
        _mm256_andnot_si256(c, y));

    *zmin = _mm256_or_si256(
        _mm256_andnot_si256(c, x),
        _mm256_and_si256(c, y));
}
```

File minmax.c – routine minmax()

y	-7	18	3	4	-1	-19	11	-10
c	0xFFFFFFFF	0x00000000	0xFFFFFFFF	0x00000000	0xFFFFFFFF	0xFFFFFFFF	0xFFFFFFFF	0xFFFFFFFF
c andnot y	Result 2	0	18	0	4	0	0	0

Minmax

Comparisons and Masks

- minmax version 2

- For illustration purpose
- Comparisons
- Masks & Boolean operations

```
void minmax2(__m256i x, __m256i y, __m256i *zmin, __m256i *zmax)
{
    __m256i c = _mm256_cmpgt_epi32(x, y);

    *zmax = _mm256_or_si256(
        _mm256_and_si256(c, x),
        _mm256_andnot_si256(c, y));

    *zmin = _mm256_or_si256(
        _mm256_andnot_si256(c, x),
        _mm256_and_si256(c, y));
}
```

File `minmax.c` — routine `minmax()`

Result 1	-3	0	4	0	7	-7	15	2	
Result 2	0	18	0	4	0	0	0	0	
zmax: R1 or R2	Result	-3	18	4	4	7	-7	15	2

Complex

Masks, Boolean ops, Integer Multiply

- **complex**
 - sum
 - $(a+ib) + (c+id) == (a+c) + i(b+d)$
 - product
 - $(a+ib) * (c+id) == (a*c - b*d) + i(a*d + b*c)$

Complex

Masks, Boolean ops, Integer Multiply

- **complex**
 - sum
 - product

```
void cmplx_sum(__m256i x, __m256i y, __m256i *z)
{
    *z = _mm256_add_epi32(x, y);
}

void cmplx_prod(__m256i x, __m256i y, __m256i *z)
{
    __m256i factor = _mm256_set_epi32(-1, 1, -1, 1, -1, 1, -1, 1);
    __m256i real_mask = _mm256_set_epi32(0, 0xffffffff, 0, 0xffffffff, 0, 0xffffffff, 0, 0xffffffff);

    __m256i real = _mm256_mullo_epi32(x, y);
    real = _mm256_mullo_epi32(factor, real);
    __m256i real_shuf = _mm256_shuffle_epi32(real, _MM_SHUFFLE(2, 3, 0, 1));
    real = _mm256_add_epi32(real, real_shuf);
    real = _mm256_and_si256(real_mask, real);

    __m256i y_shuf = _mm256_shuffle_epi32(y, _MM_SHUFFLE(2, 3, 0, 1));
    __m256i imag = _mm256_mullo_epi32(x, y_shuf);
    __m256i imag_shuf = _mm256_shuffle_epi32(imag, _MM_SHUFFLE(2, 3, 0, 1));
    imag = _mm256_add_epi32(imag, imag_shuf);
    imag = _mm256_andnot_si256(real_mask, imag);

    *z = _mm256_or_si256(real, imag);
}
```

File [complex.c](#) – routines `cmplx_sum()` & `cmplx_product()`

Complex

Masks, Boolean ops, Integer Multiply

- **complex**
 - **sum**
 - product

```
void cmplx_sum(__m256i x, __m256i y, __m256i *z)
{
    *z = _mm256_add_epi32(x, y);
}

void cmplx_prod(__m256i x, __m256i y, __m256i *z)
{
    __m256i factor = _mm256_set_epi32(-1, 1, -1, 1, -1, 1, -1, 1);
    __m256i real_mask = _mm256_set_epi32(0, 0xffffffff, 0, 0xffffffff, 0, 0xffffffff, 0, 0xffffffff);

    __m256i real = _mm256_mullo_epi32(x, y);
    real = _mm256_mullo_epi32(factor, real);
    __m256i real_shuf = _mm256_shuffle_epi32(real, _MM_SHUFFLE(2, 3, 0, 1));
    real = _mm256_add_epi32(real, real_shuf);
    real = _mm256_and_si256(real_mask, real);

    __m256i y_shuf = _mm256_shuffle_epi32(y, _MM_SHUFFLE(2, 3, 0, 1));
    __m256i imag = _mm256_mullo_epi32(x, y_shuf);
    __m256i imag_shuf = _mm256_shuffle_epi32(imag, _MM_SHUFFLE(2, 3, 0, 1));
    imag = _mm256_add_epi32(imag, imag_shuf);
    imag = _mm256_andnot_si256(real_mask, imag);

    *z = _mm256_or_si256(real, imag);
}
```

$$(a+ib) + (c+id) == (a+c) + i(b+d)$$

File `complex.c` – routines `cmplx_sum()` & `cmplx_product()`

Complex

Masks, Boolean ops, Integer Multiply

- **complex**
 - sum
 - **product**

```
void cmplx_sum(__m256i x, __m256i y, __m256i *z)
{
    *z = _mm256_add_epi32(x, y);
}

void cmplx_prod(__m256i x, __m256i y, __m256i *z)
{
    __m256i factor = _mm256_set_epi32(-1, 1, -1, 1, -1, 1, -1, 1);
    __m256i real_mask = _mm256_set_epi32(0, 0xffffffff, 0, 0xffffffff, 0, 0xffffffff, 0, 0xffffffff);

    __m256i real = _mm256_mullo_epi32(x, y);
    real = _mm256_mullo_epi32(factor, real);
    __m256i real_shuf = _mm256_shuffle_epi32(real, _MM_SHUFFLE(2, 3, 0, 1));
    real = _mm256_add_epi32(real, real_shuf);
    real = _mm256_and_si256(real_mask, real);

    __m256i y_shuf = _mm256_shuffle_epi32(y, _MM_SHUFFLE(2, 3, 0, 1));
    __m256i imag = _mm256_mullo_epi32(x, y_shuf);
    __m256i imag_shuf = _mm256_shuffle_epi32(imag, _MM_SHUFFLE(2, 3, 0, 1));
    imag = _mm256_add_epi32(imag, imag_shuf);
    imag = _mm256_andnot_si256(real_mask, imag);

    *z = _mm256_or_si256(real, imag);
}
```

File [complex.c](#) – routines `cmplx_sum()` & `cmplx_product()`

Complex

Masks, Boolean ops, Integer Multiply

$$(a+ib) * (c+id) == (a*c - b*d) + i(a*d + b*c)$$

- **complex**
 - sum
 - **product**

```
void cmplx_prod(__m256i x, __m256i y, __m256i *z)
{
    __m256i factor = _mm256_set_epi32(-1, 1, -1, 1, -1, 1, -1, 1);
    __m256i real_mask = _mm256_set_epi32(0, 0xffffffff, 0, 0xffffffff, 0, 0xffffffff, 0, 0xffffffff);

    __m256i real = _mm256_mullo_epi32(x, y);
    real = _mm256_mullo_epi32(factor, real);
    __m256i real_shuf = _mm256_shuffle_epi32(real, _MM_SHUFFLE(2, 3, 0, 1));
    real = _mm256_add_epi32(real, real_shuf);
    real = _mm256_and_si256(real_mask, real);

    __m256i y_shuf = _mm256_shuffle_epi32(y, _MM_SHUFFLE(2, 3, 0, 1));
    __m256i imag = _mm256_mullo_epi32(x, y_shuf);
    __m256i imag_shuf = _mm256_shuffle_epi32(imag, _MM_SHUFFLE(2, 3, 0, 1));
    imag = _mm256_add_epi32(imag, imag_shuf);
    imag = _mm256_andnot_si256(real_mask, imag);

    *z = _mm256_or_si256(real, imag);
}
```

File [complex.c](#) – routines `cmplx_product()`

Complex

Masks, Boolean ops, Integer Multiply

$$(a+ib) * (c+id) == (a*c - b*d) + i(a*d + b*c)$$

- **complex**
 - sum
 - **product**

```
void cmplx_prod(__m256i x, __m256i y, __m256i *z)
{
    __m256i factor = _mm256_set_epi32(-1, 1, -1, 1, -1, 1, -1, 1);
    __m256i real_mask = _mm256_set_epi32(0, 0xffffffff, 0, 0xffffffff, 0, 0xffffffff, 0, 0xffffffff);

    __m256i real = _mm256_mullo_epi32(x, y);
    real = _mm256_mullo_epi32(factor, real);
    __m256i real_shuf = _mm256_shuffle_epi32(real, _MM_SHUFFLE(2, 3, 0, 1));
    real = _mm256_add_epi32(real, real_shuf);
    real = _mm256_and_si256(real_mask, real);

    __m256i y_shuf = _mm256_shuffle_epi32(y, _MM_SHUFFLE(2, 3, 0, 1));
    __m256i imag = _mm256_mullo_epi32(x, y_shuf);
    __m256i imag_shuf = _mm256_shuffle_epi32(imag, _MM_SHUFFLE(2, 3, 0, 1));
    imag = _mm256_add_epi32(imag, imag_shuf);
    imag = _mm256_andnot_si256(real_mask, imag);

    *z = _mm256_or_si256(real, imag);
}
```

File [complex.c](#) – routines `cmplx_product()`

Complex

Masks, Boolean ops, Integer Multiply

$$(a+ib) * (c+id) == (a*c - b*d) + i(a*d + b*c)$$

- **complex**
 - sum
 - **product**

```
void cmplx_prod(__m256i x, __m256i y, __m256i *z)
{
    __m256i factor = _mm256_set_epi32(-1, 1, -1, 1, -1, 1, -1, 1);
    __m256i real_mask = _mm256_set_epi32(0, 0xffffffff, 0, 0xffffffff, 0, 0xffffffff, 0, 0xffffffff);

    __m256i real = _mm256_mullo_epi32(x, y);
    real = _mm256_mullo_epi32(factor, real);
    __m256i real_shuf = _mm256_shuffle_epi32(real, _MM_SHUFFLE(2, 3, 0, 1));
    real = _mm256_add_epi32(real, real_shuf);
    real = _mm256_and_si256(real_mask, real);

    __m256i y_shuf = _mm256_shuffle_epi32(y, _MM_SHUFFLE(2, 3, 0, 1));
    __m256i imag = _mm256_mullo_epi32(x, y_shuf);
    __m256i imag_shuf = _mm256_shuffle_epi32(imag, _MM_SHUFFLE(2, 3, 0, 1));
    imag = _mm256_add_epi32(imag, imag_shuf);
    imag = _mm256_andnot_si256(real_mask, imag);

    *z = _mm256_or_si256(real, imag);
}
```

File [complex.c](#) – routines `cmplx_product()`

Complex

Masks, Boolean ops, Integer Multiply

$$(a+ib) * (c+id) == (a*c - b*d) + i(a*d + b*c)$$

- **complex**
 - sum
 - **product**

```
void cmplx_prod(__m256i x, __m256i y, __m256i *z)
{
    __m256i factor = _mm256_set_epi32(-1, 1, -1, 1, -1, 1, -1, 1);
    __m256i real_mask = _mm256_set_epi32(0, 0xffffffff, 0, 0xffffffff, 0, 0xffffffff, 0, 0xffffffff);

    __m256i real = _mm256_mullo_epi32(x, y);
    real = _mm256_mullo_epi32(factor, real);
    __m256i real_shuf = _mm256_shuffle_epi32(real, _MM_SHUFFLE(2, 3, 0, 1));
    real = _mm256_add_epi32(real, real_shuf);
    real = _mm256_and_si256(real_mask, real);

    __m256i y_shuf = _mm256_shuffle_epi32(y, _MM_SHUFFLE(2, 3, 0, 1));
    __m256i imag = _mm256_mullo_epi32(x, y_shuf);
    __m256i imag_shuf = _mm256_shuffle_epi32(imag, _MM_SHUFFLE(2, 3, 0, 1));
    imag = _mm256_add_epi32(imag, imag_shuf);
    imag = _mm256_andnot_si256(real_mask, imag);

    *z = _mm256_or_si256(real, imag);
}
```

File [complex.c](#) – routines `cmplx_product()`

Complex

Masks, Boolean ops, Integer Multiply

$$(a+ib) * (c+id) == (a*c - b*d) + i(a*d + b*c)$$

- **complex**
 - sum
 - **product**

```
void cmplx_prod(__m256i x, __m256i y, __m256i *z)
{
    __m256i factor = _mm256_set_epi32(-1, 1, -1, 1, -1, 1, -1, 1);
    __m256i real_mask = _mm256_set_epi32(0, 0xffffffff, 0, 0xffffffff, 0, 0xffffffff, 0, 0xffffffff);

    __m256i real = _mm256_mullo_epi32(x, y);
    real = _mm256_mullo_epi32(factor, real);
    __m256i real_shuf = _mm256_shuffle_epi32(real, _MM_SHUFFLE(2, 3, 0, 1));
    real = _mm256_add_epi32(real, real_shuf);
    real = _mm256_and_si256(real_mask, real);

    __m256i y_shuf = _mm256_shuffle_epi32(y, _MM_SHUFFLE(2, 3, 0, 1));
    __m256i imag = _mm256_mullo_epi32(x, y_shuf);
    __m256i imag_shuf = _mm256_shuffle_epi32(imag, _MM_SHUFFLE(2, 3, 0, 1));
    imag = _mm256_add_epi32(imag, imag_shuf);
    imag = _mm256_andnot_si256(real_mask, imag);

    *z = _mm256_or_si256(real, imag);
}
```

File [complex.c](#) – routines `cmplx_product()`

Examples of Common SIMD Instructions – AVX2

Masked Load / Masked Store

- **Masked Load**
 - Load data from memory into a SIMD register
 - `_mm256_maskload_*`
- **Masked Store**
 - Store data from a SIMD register into memory
 - `_mm256_maskstore_*`

Examples of Common SIMD Instructions – AVX2

Masked Load / Masked Store

- **Masked Load**

- Load data from memory into a SIMD register
 - `_mm256_maskload_*`

Mem	-3		-1		4		-12		7		-7		15		2		
Mask	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	bit number
	1	?	0	?	0	?	1	?	0	?	1	?	1	?	1	?	bit value

- **Masked Store**

- Store data from a SIMD register into memory
 - `_mm256_maskstore_*`

Examples of Common SIMD Instructions – AVX2

Masked Load / Masked Store

- **Masked Load**

- Load data from memory into a SIMD register
 - `_mm256_maskload_*`

Mem	-3		-1		4		-12		7		-7		15		2		
Mask	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	bit number
	1	?	0	?	0	?	1	?	0	?	1	?	1	?	1	?	bit value
Reg	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	

- **Masked Store**

- Store data from a SIMD register into memory
 - `_mm256_maskstore_*`

Examples of Common SIMD Instructions – AVX2

Masked Load / Masked Store

- **Masked Load**

- Load data from memory into a SIMD register
 - `_mm256_maskload_*`

Mem	-3	-1	4	-12	7	-7	15	2									
Mask	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	bit number
Reg	-3	0	0	-12	0	-7	15	2	bit value								

- **Masked Store**

- Store data from a SIMD register into memory
 - `_mm256_maskstore_*`

Examples of Common SIMD Instructions – AVX2

Masked Load / Masked Store

- **Masked Load**
 - Load data from memory into a SIMD register
 - `_mm256_maskload_*`
- **Masked Store**
 - Store data from a SIMD register into memory
 - `_mm256_maskstore_*`

Reg	-16				8				15				-18				-6				13				1		14	
Mask	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	31	30 ... 0	bit number	
	1	?	0	?	0	?	1	?	0	?	1	?	1	?	1	?	1	?	1	?	1	?	1	?	1	?	bit value	

Examples of Common SIMD Instructions – AVX2

Masked Load / Masked Store

- **Masked Load**
 - Load data from memory into a SIMD register
 - `_mm256_maskload_*`
- **Masked Store**
 - Store data from a SIMD register into memory
 - `_mm256_maskstore_*`

Reg	-16	8	15	-18	-6	13	1	14							
Mask	31	30...0	31	30...0	31	30...0	31	30...0	31	30...0	31	30...0	31	30...0	bit number
Mem	1	?	0	?	0	?	1	?	0	?	1	?	1	?	bit value

Examples of Common SIMD Instructions – AVX2

Masked Load / Masked Store

- **Masked Load**
 - Load data from memory into a SIMD register
 - `_mm256_maskload_*`
- **Masked Store**
 - Store data from a SIMD register into memory
 - `_mm256_maskstore_*`

Reg	-16	8	15	-18	-6	13	1	14							
Mask	31	30...0	31	30...0	31	30...0	31	30...0	31	30...0	31	30...0	31	30...0	bit number
Mem	1	?	0	?	0	?	1	?	0	?	1	?	1	?	bit value

TD

Introduction to using *intrinsics* routines for the Intel AVX2 instruction set, part II

- Assignment, source codes, additional slides
 - Moodle ENSEIRB, IT390 course

